

# Deep Neural Networks And Where to Find Them

Artem Korenev, Nikita Gryaznov

Skoltech ISP, 2018

# Welcome

- Hi everybody!
- Two TAs: Artem and Nikita (2<sup>nd</sup> year MSc students)
- 6 days in-class (~20 hours)
- 40 hours for the course in total
- Other time for HW and self education
- If you don't get something – you are very welcome to ask a question right away
- Don't be shy and come after lecture if you need

# What the Course is About

- Introduction in Deep Neural Networks and how to use them
- This is mostly an introductory course (your first starting point in DL)
- We aim on people outside IT track
- Bias towards practice instead of theory
- We will use Keras library (not Tensorflow, not Theano, not PyTorch)
- Currently, we are not going to even briefly touch things like GAN, RNN, Reinforcement Learning, Segmentation and Detection
- If you are an advanced student and still want to learn DL further please come forward and we will discuss what you can do individually

# Syllabus\*

Days:

1. ML introduction. Logistic regression. Neural networks approach.
  2. Matrix representation of neural networks. Details of NNs optimization.
  3. Activation functions. Loss functions. Regression problem.
  4. Different optimizers. Regularization.
  5. Convolutional Neural Networks basics
  6. Practical tricks and consultation.
- + Daily seminars and exercises climaxing with a final homework assignment

*\* Subject to change depending on the pace*

# Requirements for the Course

- Basic proficiency in Python
- Simple math knowledge (matrices, derivatives)
- Having a laptop
  
- None of this is absolutely mandatory but will help you understand many things
- Let's have a quiz everybody: <https://goo.gl/ZTXf3v>

# Why Bother?

A Couple of Words About Machine Learning

# Machine Learning in 2018

- The sexiest job of the 21<sup>st</sup> century
- Creating programs from data
- Field has a lot of algorithms (kNN, decision trees, SVM, etc.)
- Neural Networks is just one of them
- NNs are growing crazy since 2012
- NNs are everywhere
- Let's see some examples...


# Image classification

- ImageNet Challenge, 2012






# Object Detection

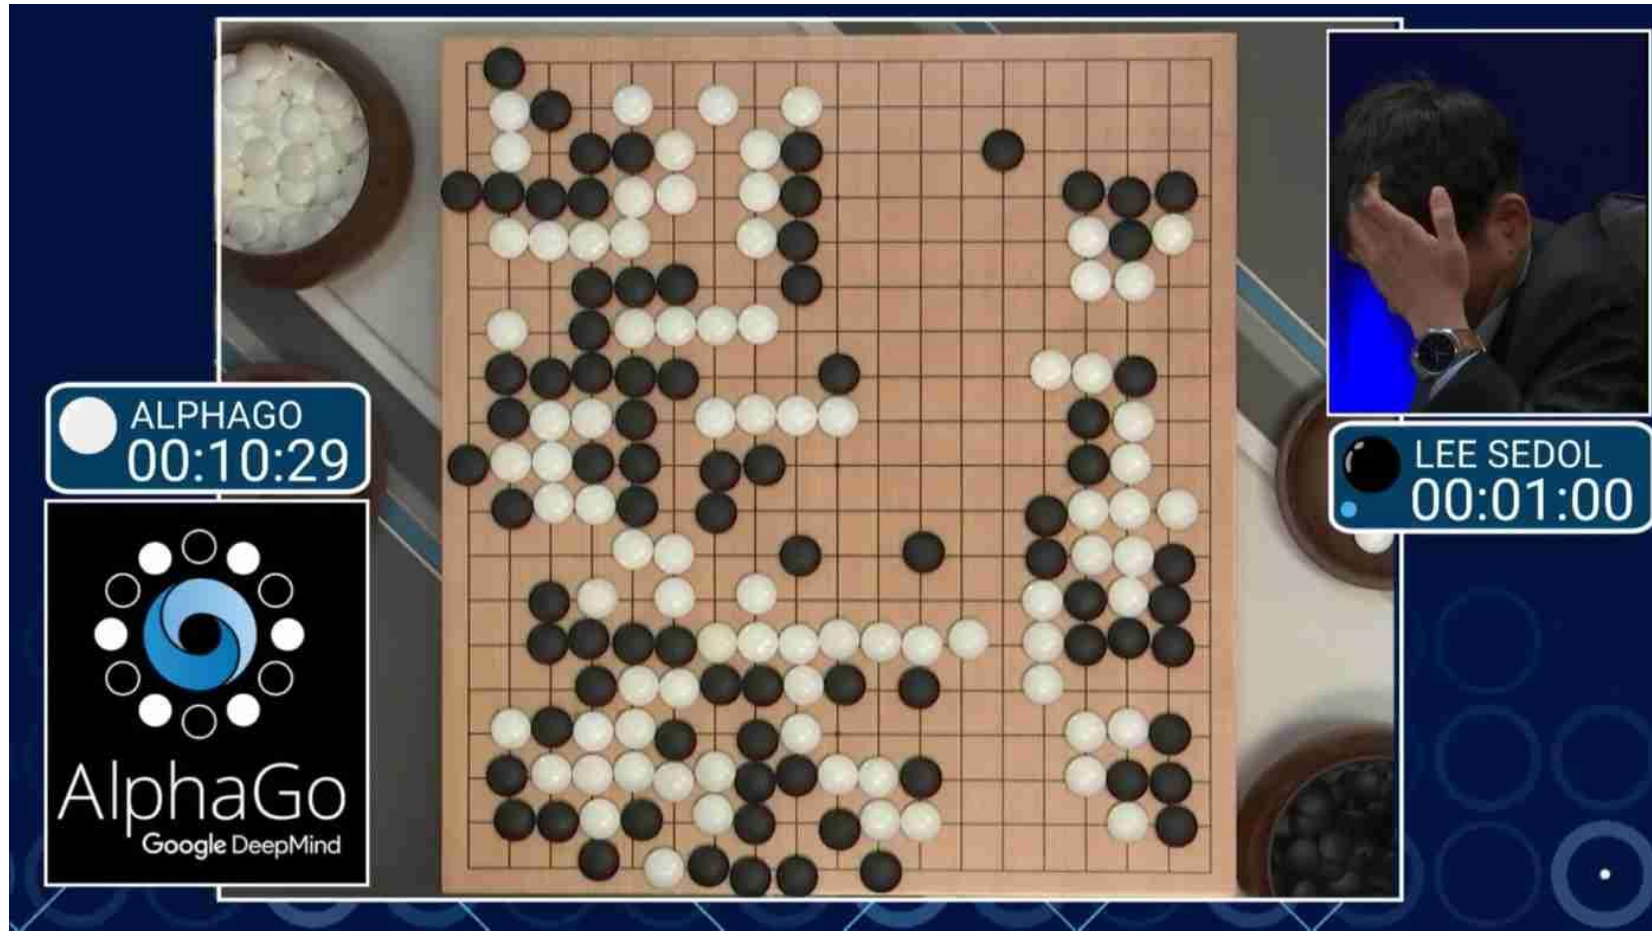


<https://www.youtube.com/watch?v=VOC3huqHrss>

# Image Segmentation

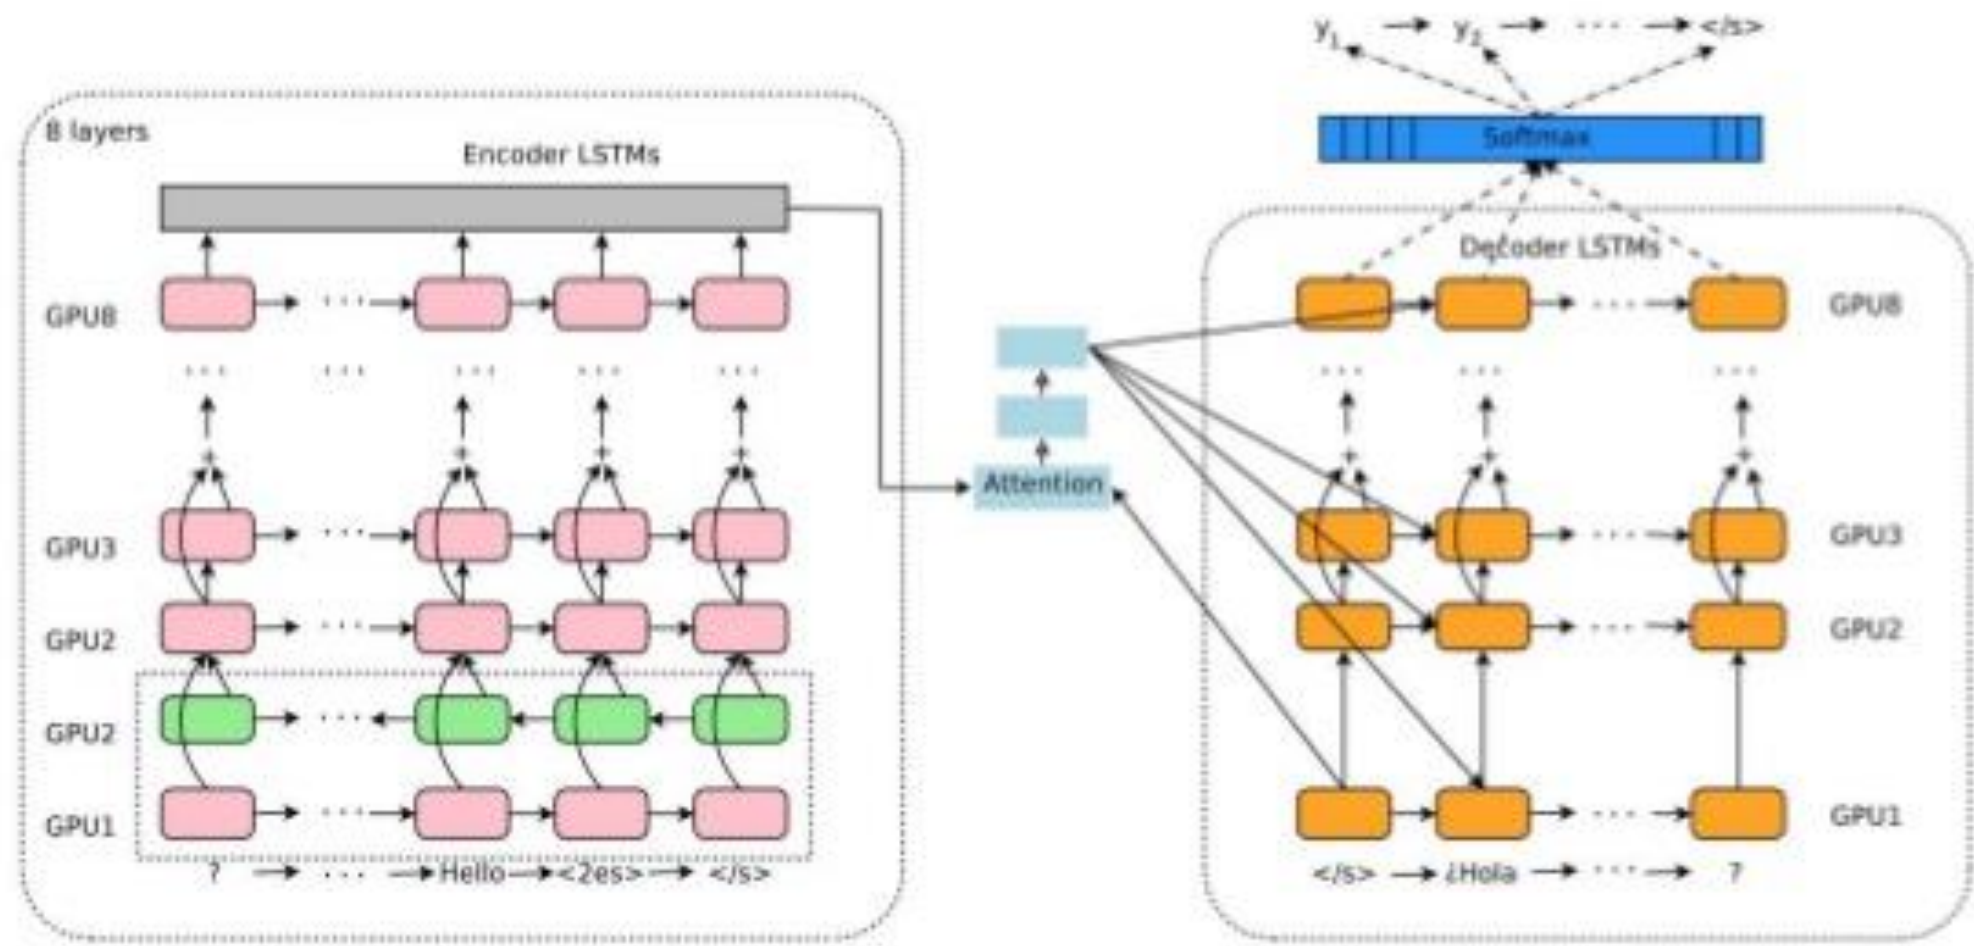


# AlphaGo





# Machine Translation



# Prisma App



# Demos

- <http://make.girls.moe/>
- <https://www.youtube.com/watch?v=9reHvktowLY>
- <https://affinelayer.com/pixsrv/>
- <https://www.youtube.com/watch?v=DgPaCWJL7XI>
- <https://google.github.io/tacotron/publications/tacotron2/index.html>
- <http://www.aiva.ai/>
- And many more

# Machine Learning Algorithms

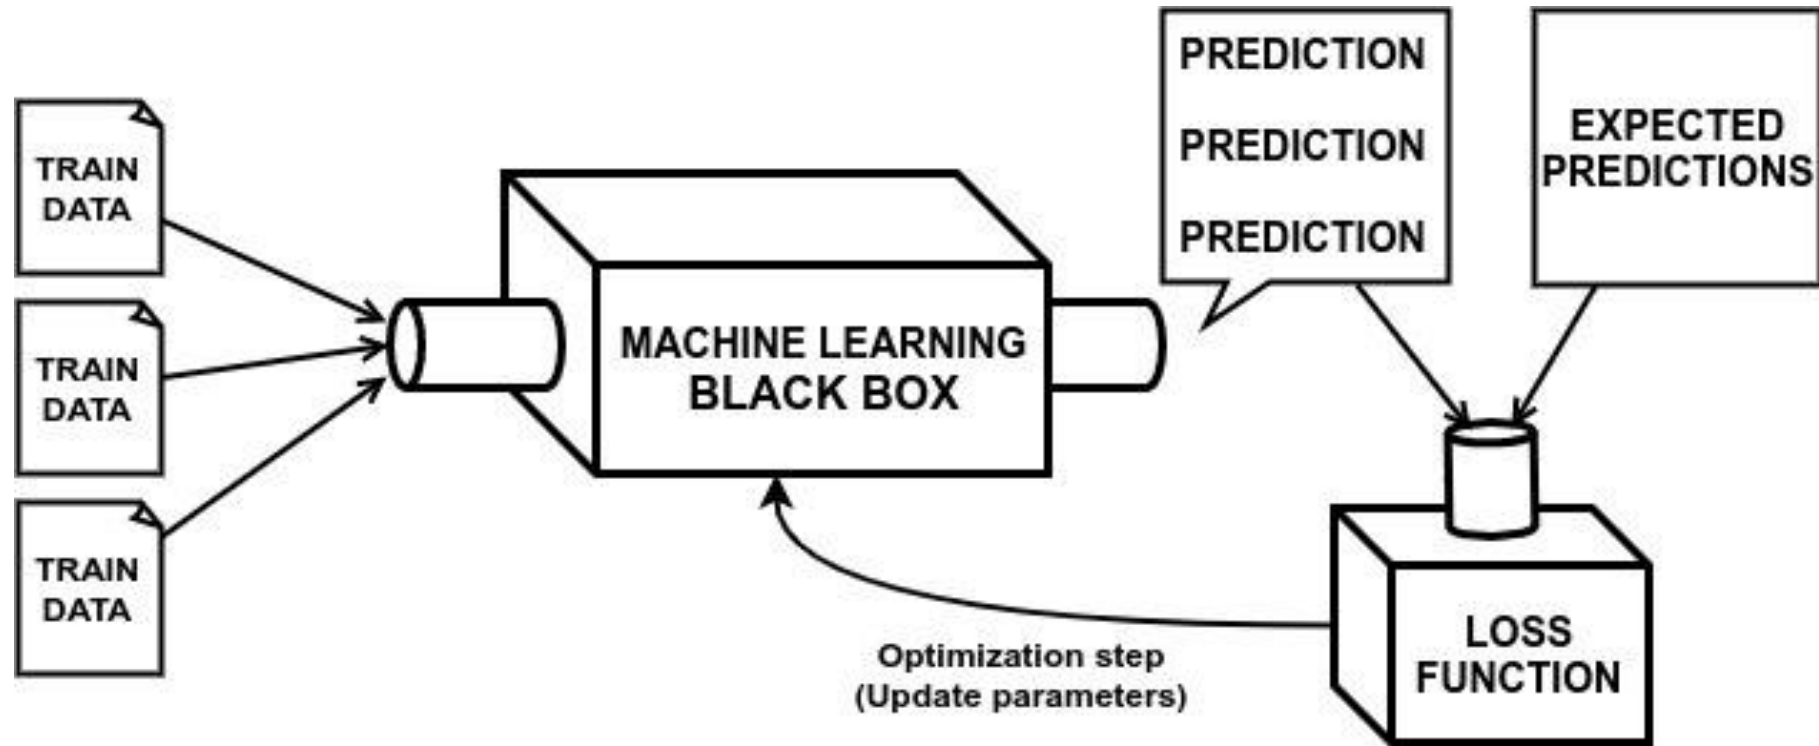
How they (usually) work

# Machine Learning Algorithms

- We create a box that has needed input and output (called *model*)
- For classification problem, for example, an image and the number of a class
- We design a function that measures badness of this black box (called *loss function*)
- Given data we try to tweak or *model* to achieve the minimal score for *loss function*

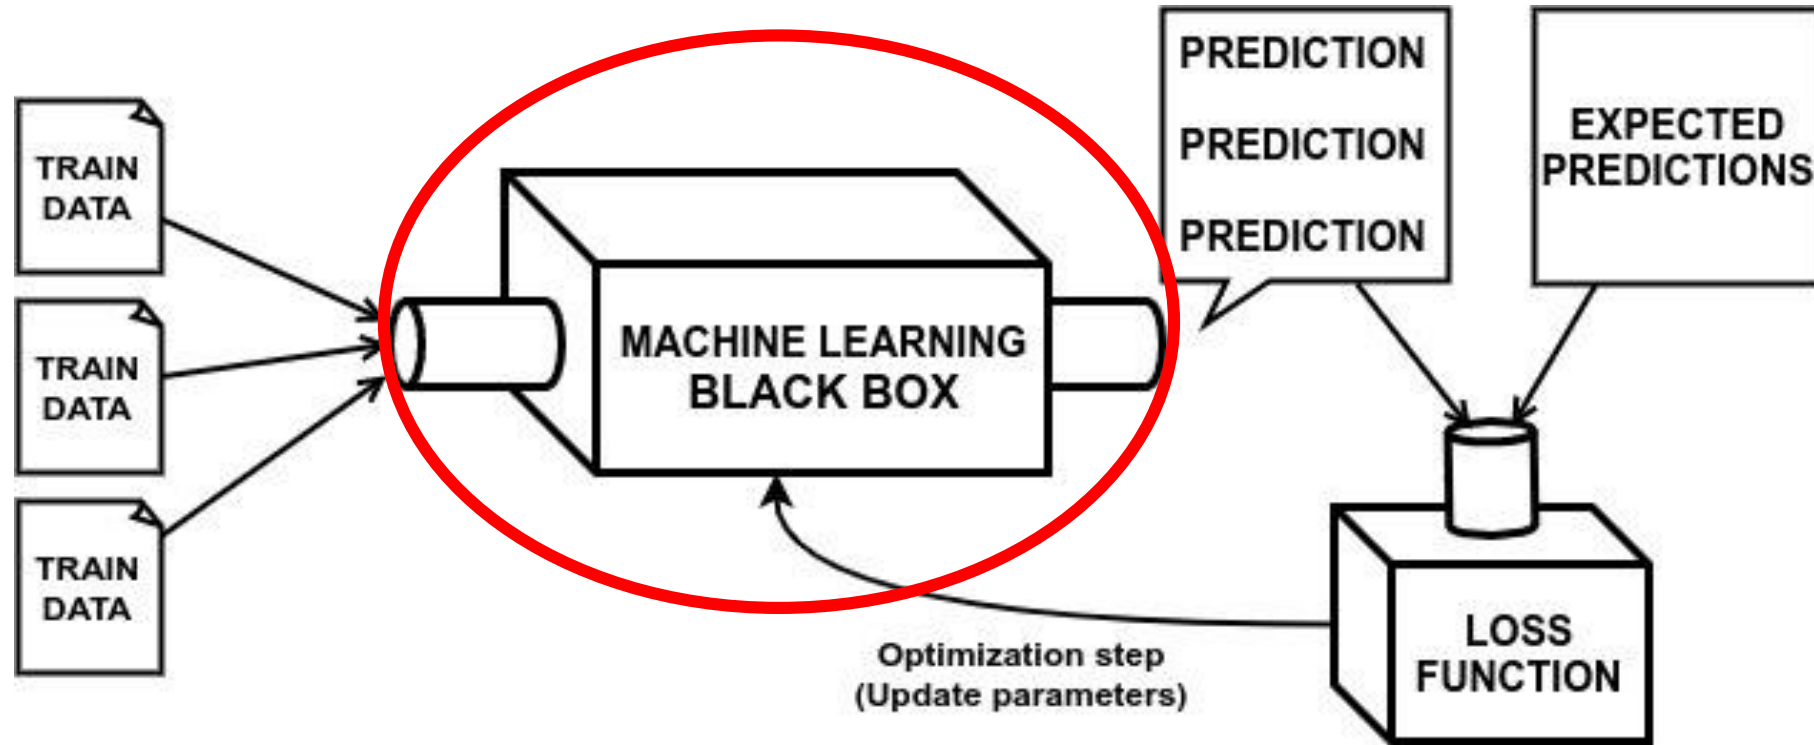


# Machine Learning Algorithms

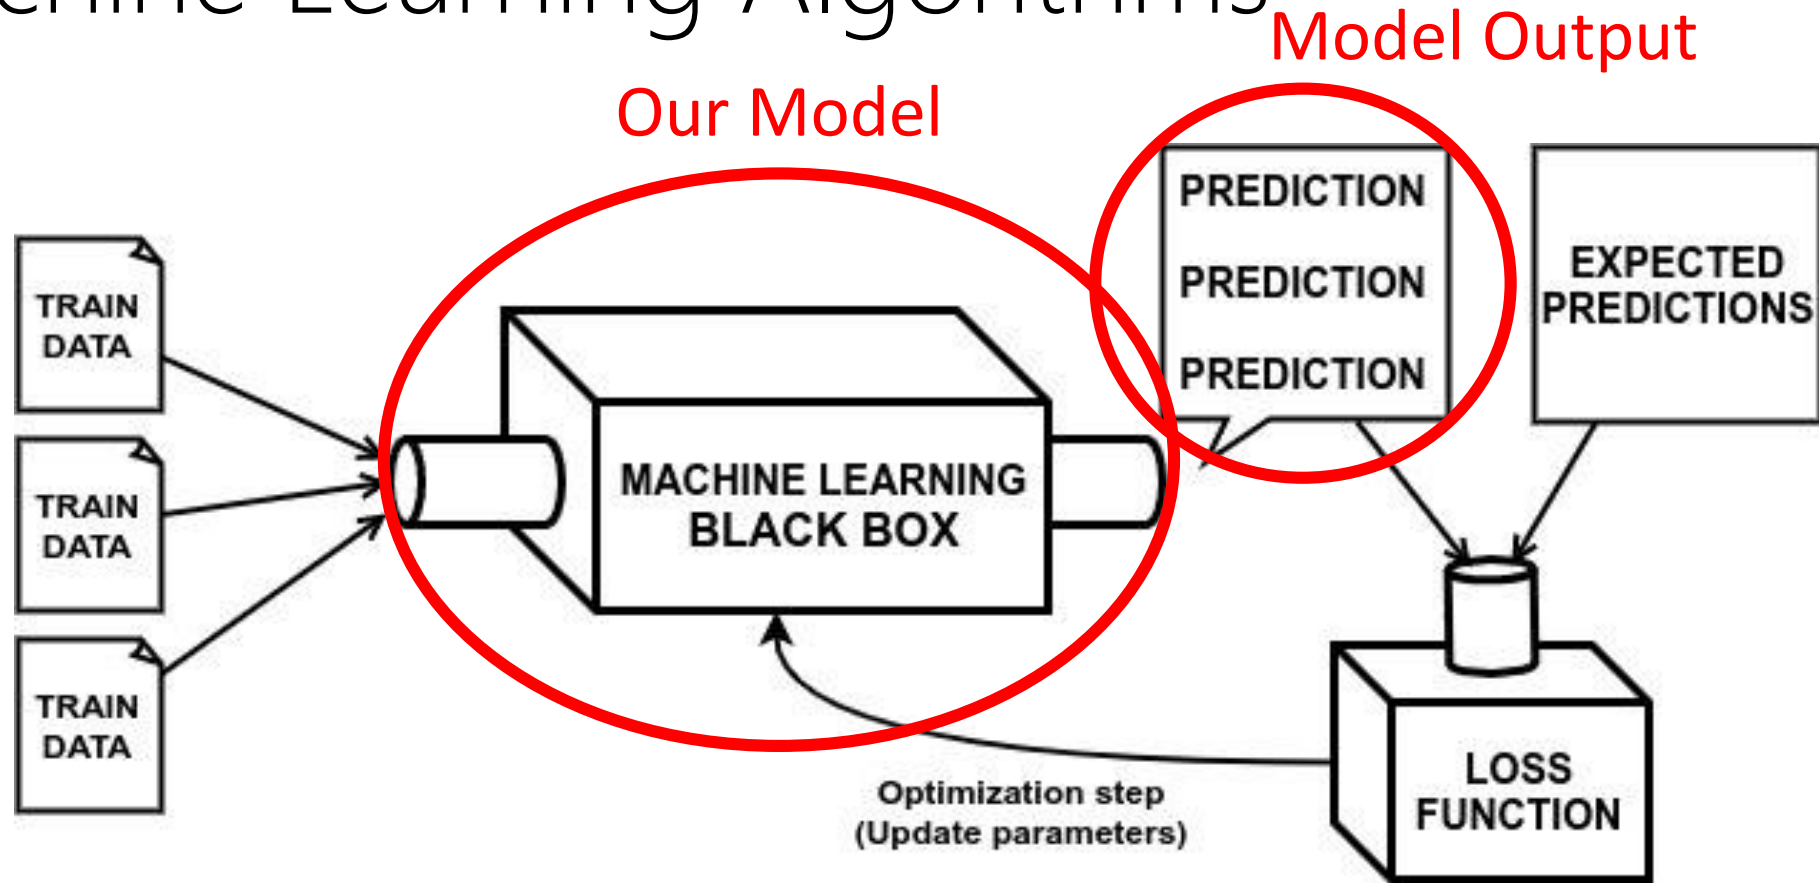


# Machine Learning Algorithms

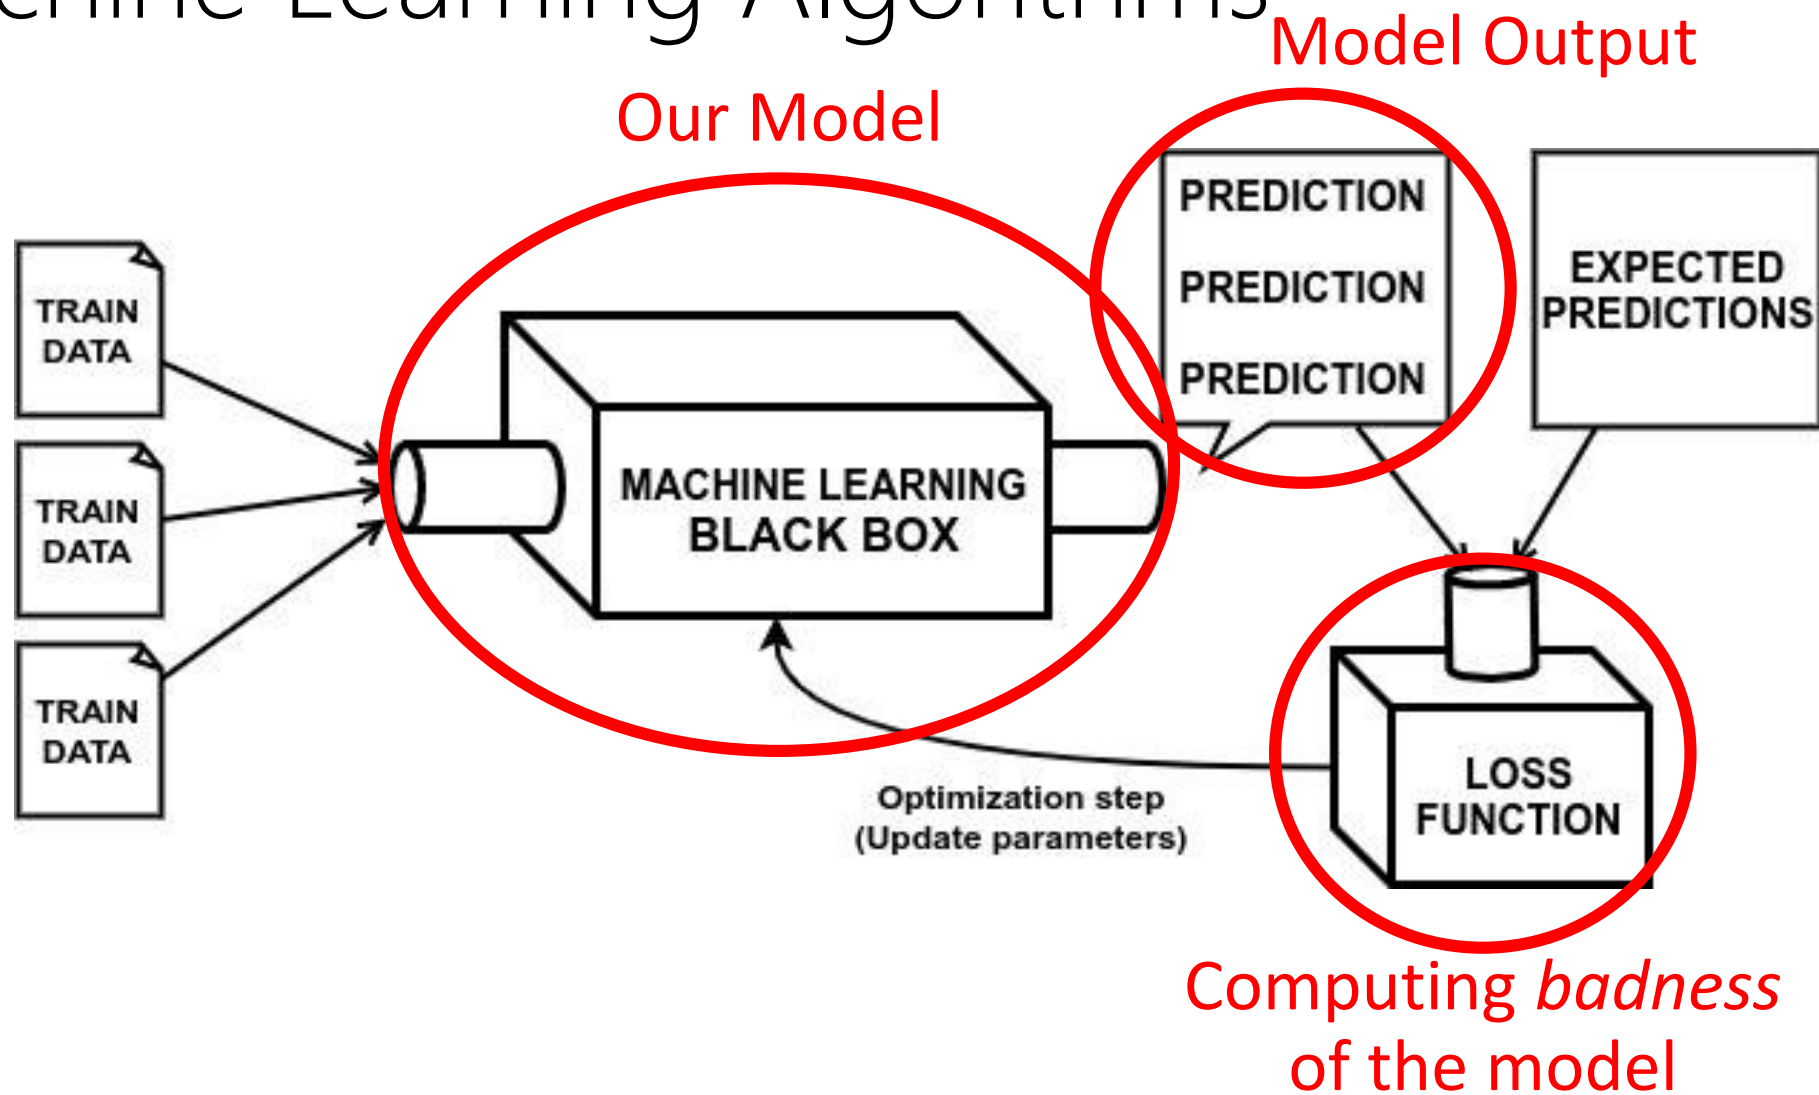
## Our Model



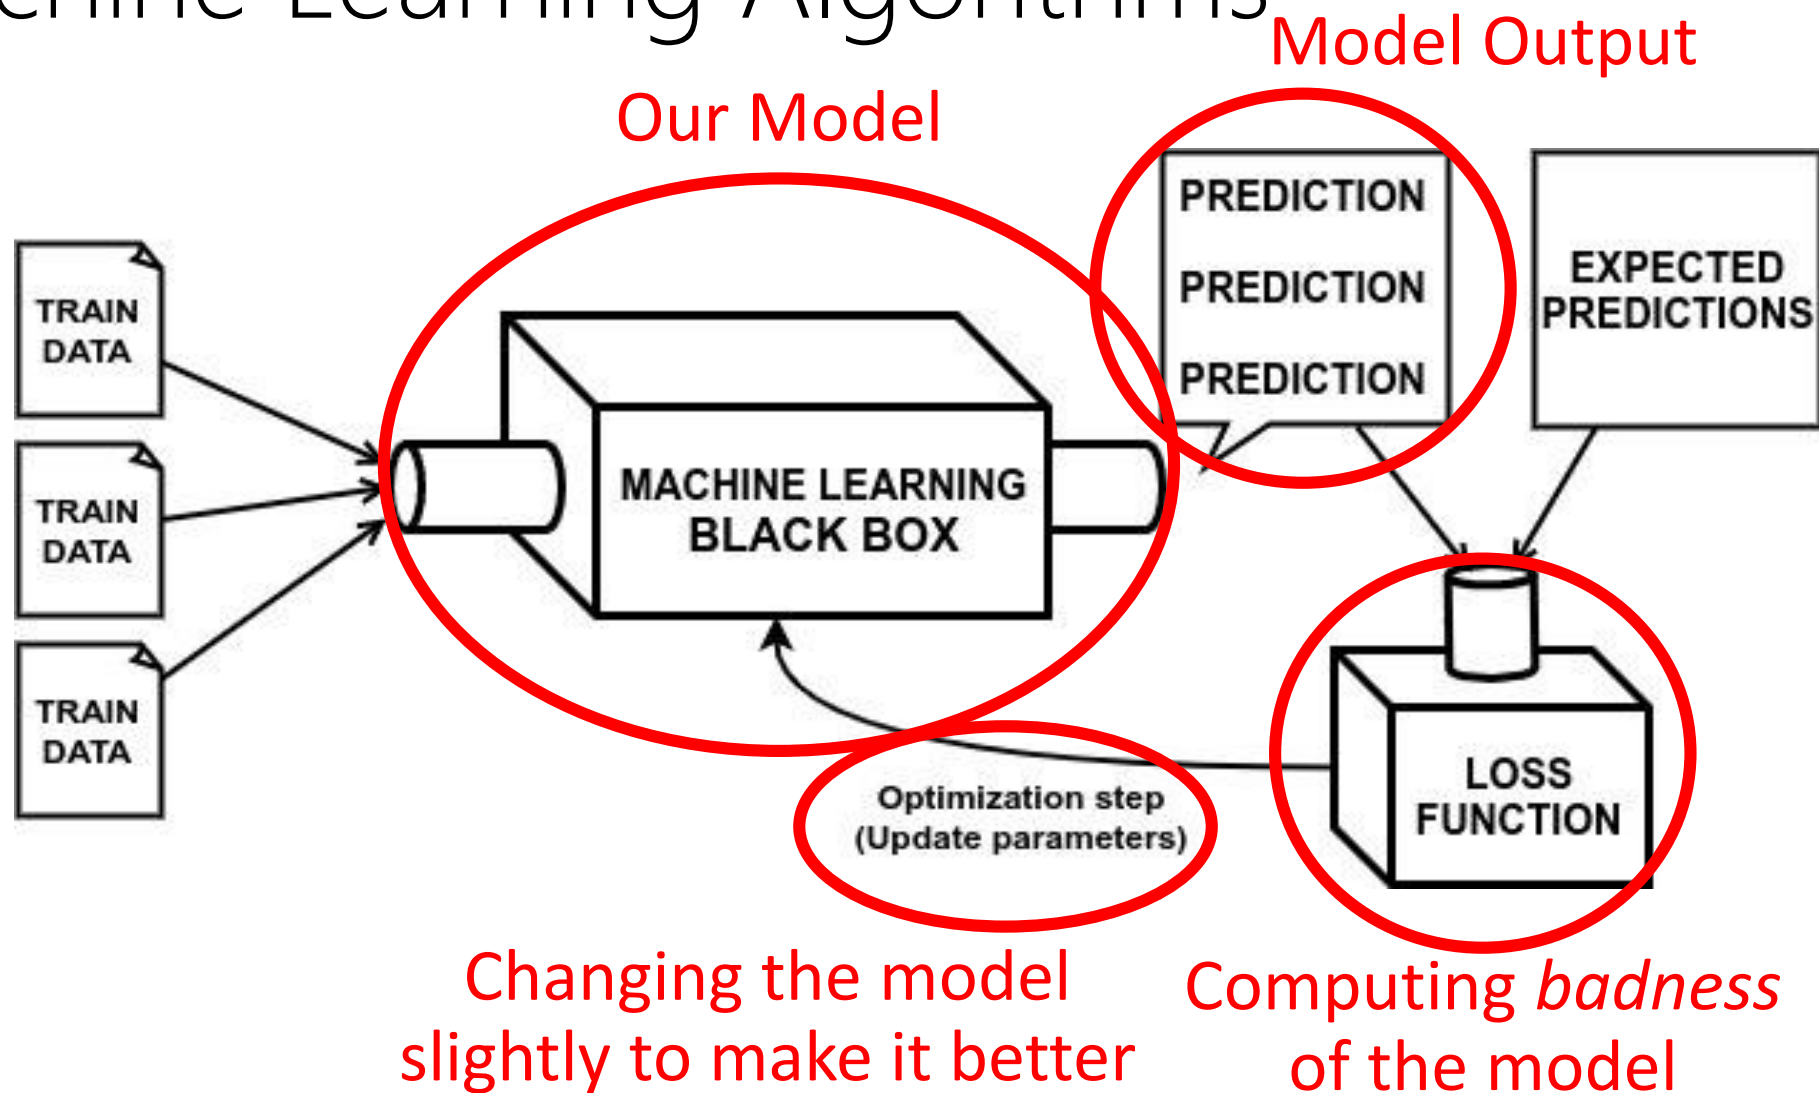
# Machine Learning Algorithms



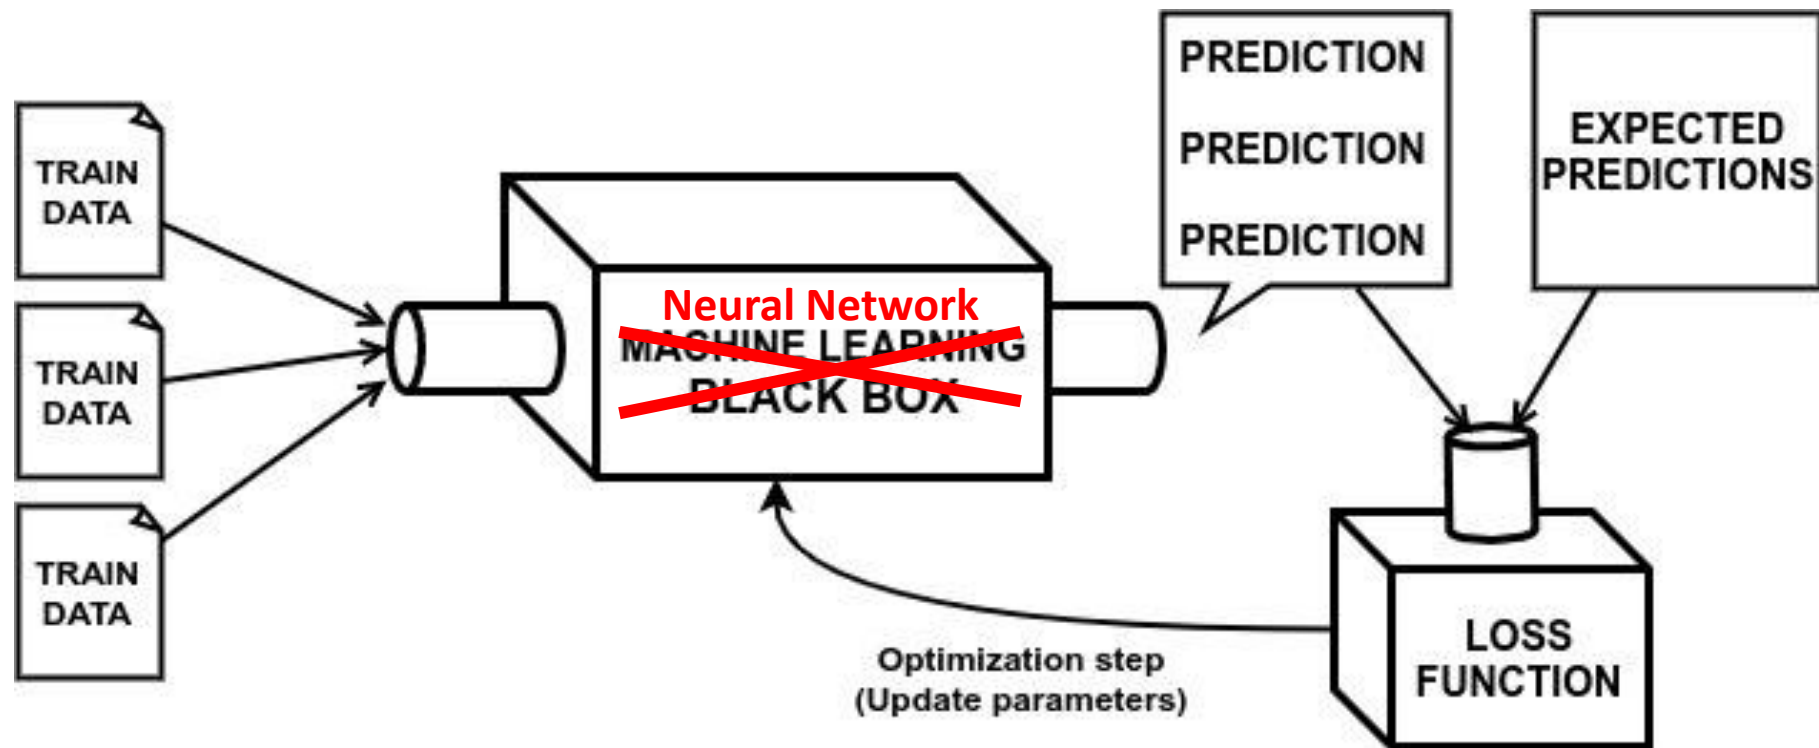
# Machine Learning Algorithms



# Machine Learning Algorithms

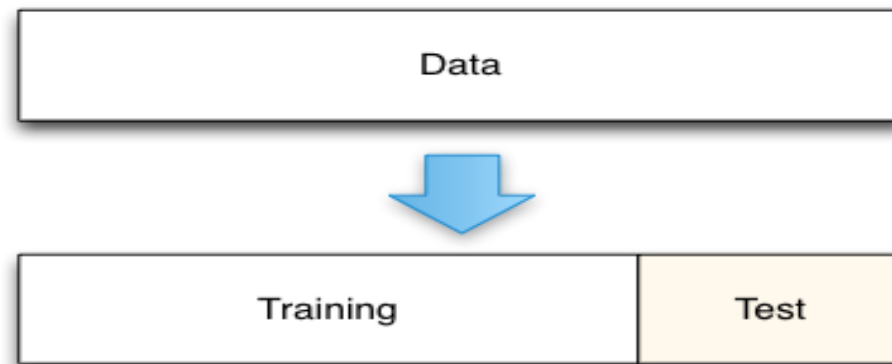


# Neural Network as a ML Black Box





# A Bit About Data

- To train a model we use data
- Models sometimes become really complex so they just memorize dataset instead of generalizing the problem
- To spot this we perform splitting data to *training* and *testing* data
- We train model using *training data* and measure score on *testing data*



# Overfitting





# What Types of Problems Can NNs Solve?

Most popular ones:


- Classification (output one of the predefined classes  $\{0, 1, \dots, n\}$ )
- Regression (output real value) (e.g. predict size of something)
- Many tasks are just combinations of these two
- There are exceptions (but we are not going that deep)

# Logistic Regression Algorithm

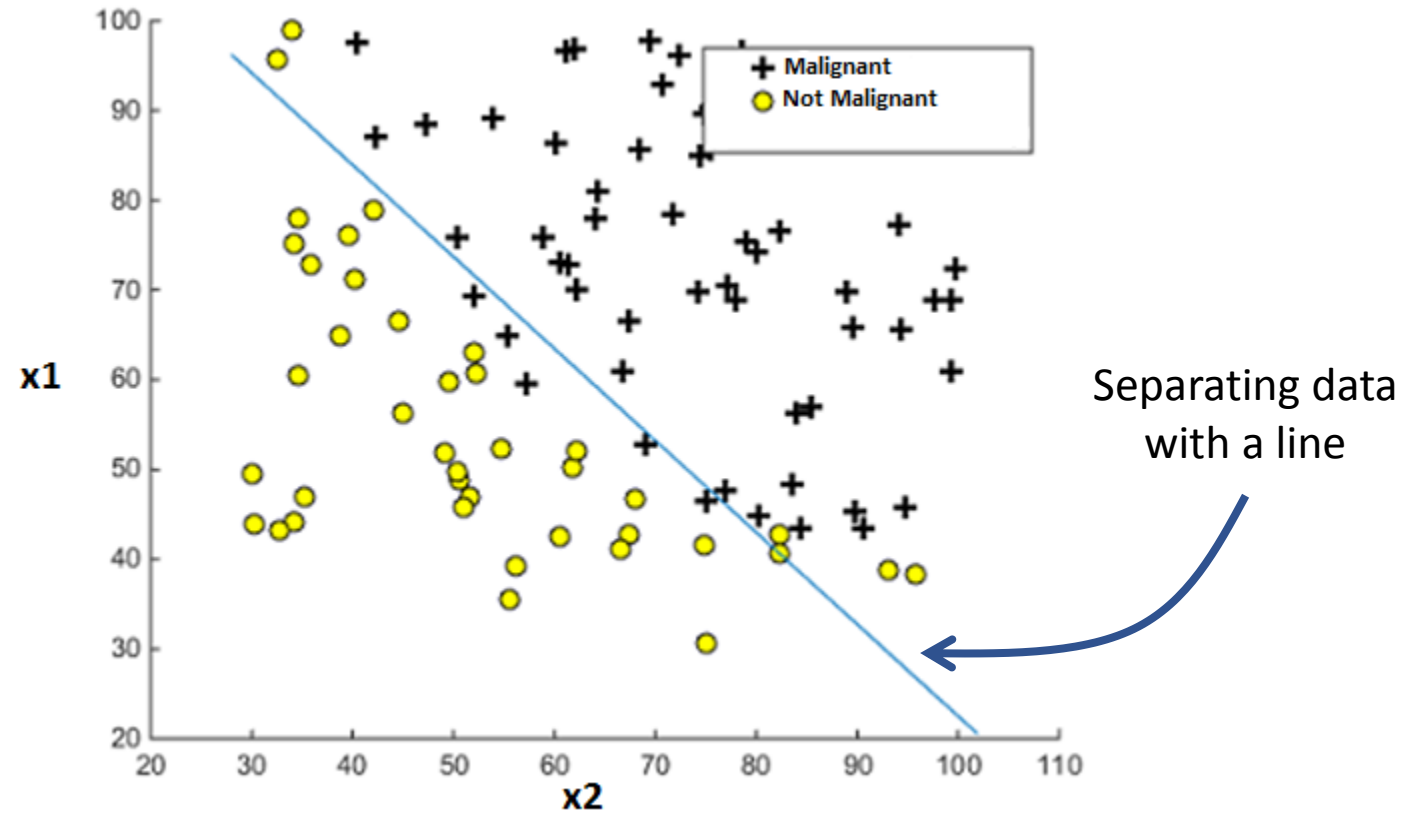
# Logistic Regression

- A simple binary classification algorithm
- Output: 0 or 1.
- Input: data with  $n$  parameters
  
- Example: prediction of tumor malignancy
- Output: *not malignant (0)* or *malignant (1)*
- Input: data with parameters: *tumor size, age, tumor age...*

# Linear Data Separation



# Linear Data Separation



# Linear Data Separation

- We create equation for the line:

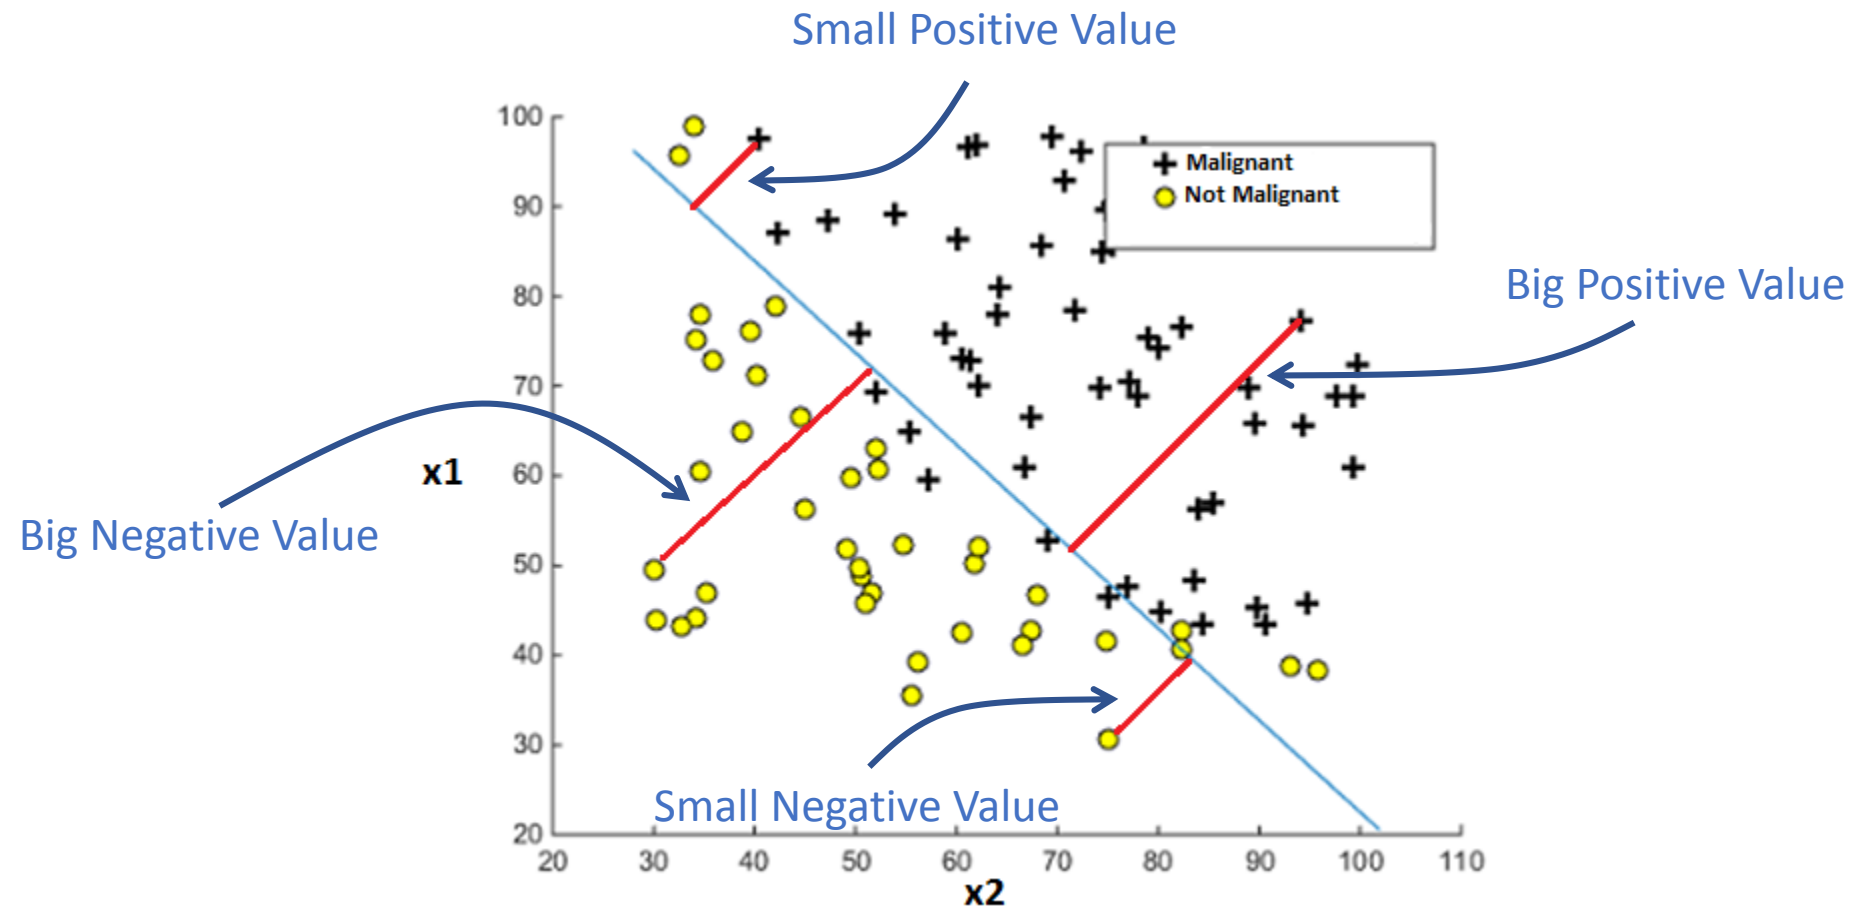
$$w_1x_1 + w_2x_2 + \cdots + w_nx_n + b = 0$$

- To measure distance from this line we can use

$$W(x) = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

- Sign of  $W(x)$  would say what class the data point has


# Linear Data Separation



# Linear Data Separation

- We can transform distance from the line to *probability* of a data point having a positive class
- We can use *sigmoid* function to do that

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



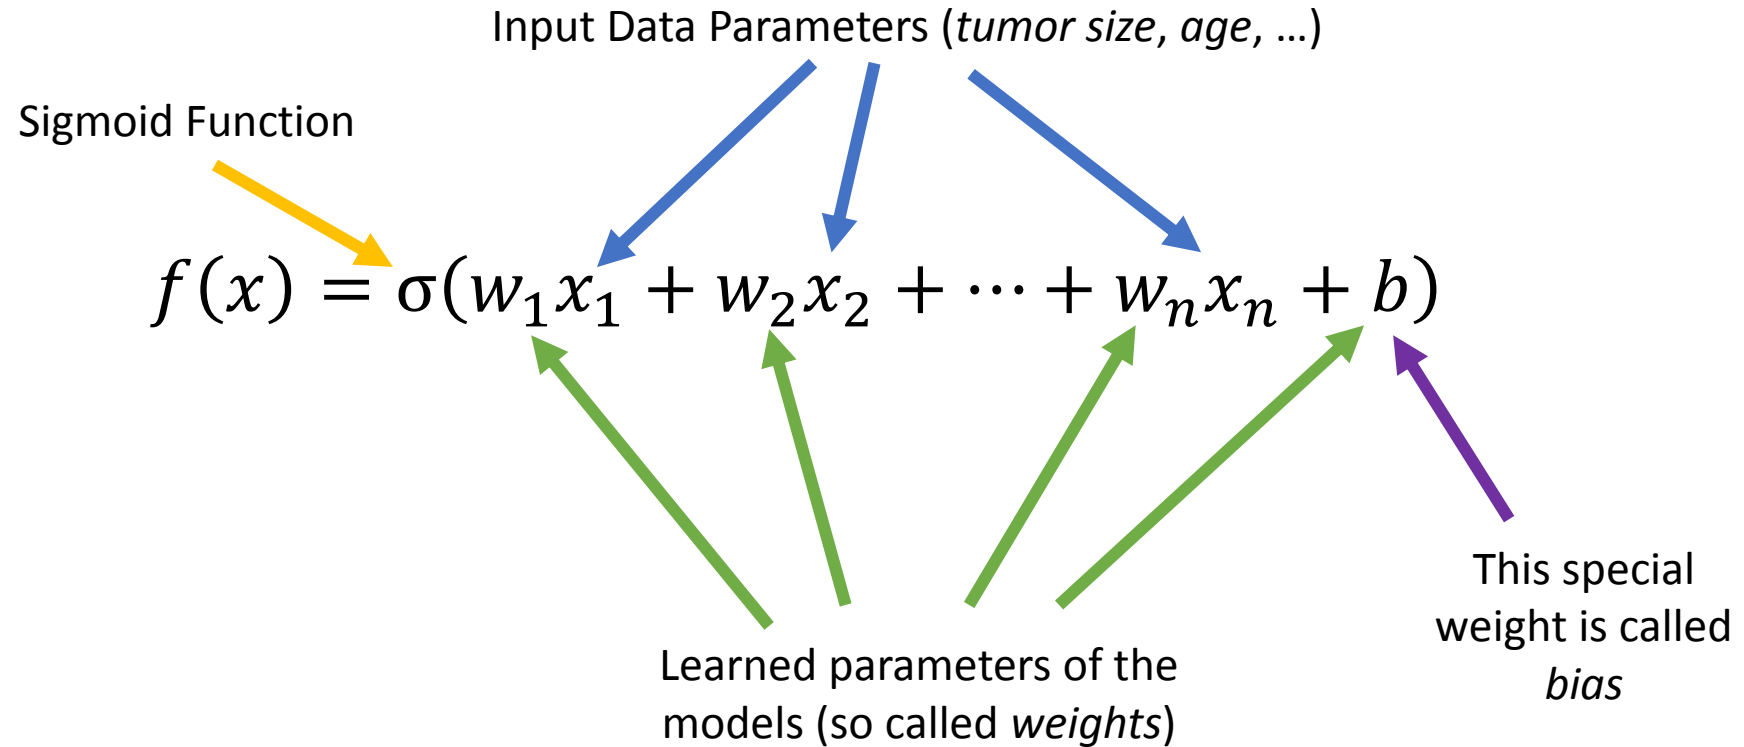


# Logistic Regression

More strict way

- Input:  $x = x_1, x_2, \dots, x_n$  (just  $n$  parameters for an object)
- Output:  $f(x) = \sigma(w_1x_1 + w_2x_2 + \dots + w_nx_n + b) \in [0; 1]$
- Sigmoid function:  $\sigma(x) = \frac{1}{1+e^{-x}}$
- For final decision  $\{0, 1\}$  we can just round the output:
  - 0 if  $f(x) < 0.5$
  - 1 if  $f(x) \geq 0.5$

# Logistic Regression – Closer Look



# How Do Measure Badness?

- We need an objective function (loss function)
- Cross Entropy:

$$L(y, f(x)) = -\frac{1}{n} \sum_i^n [y_i \log(f(x)) + (1 - y_i) \log(1 - f(x))]$$

- (Derived from Maximum Likelihood Estimation)

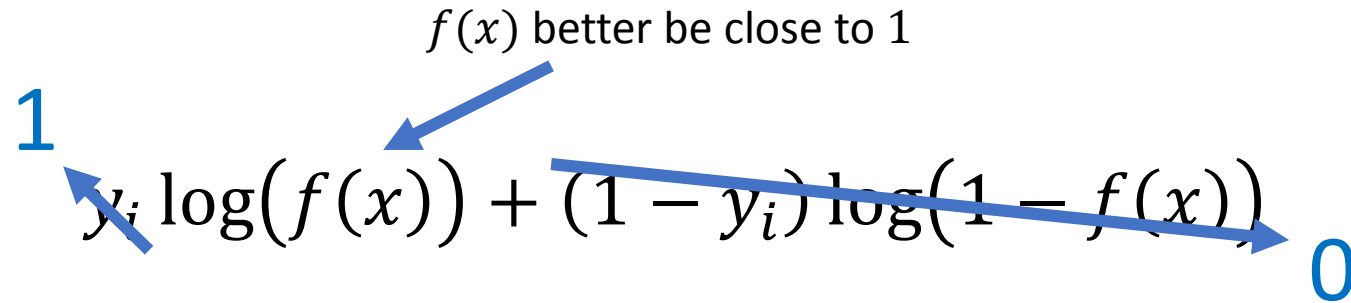
# Looking Closer on the Objective

$$L(y, f(x)) = -\frac{1}{n} \sum_i^n [y_i \log(f(x)) + (1 - y_i) \log(1 - f(x))]$$

and  $f(x) \in [0; 1]$

If  $y_i = 1$

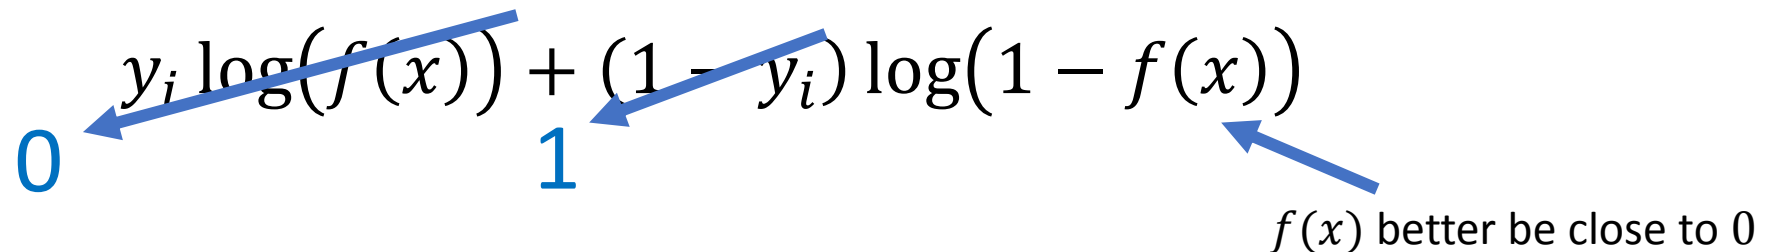
$f(x)$  better be close to 1



The diagram shows the term  $y_i \log(f(x)) + (1 - y_i) \log(1 - f(x))$  with a blue '1' to the left of the first term and a blue '0' to the right of the second term. A blue arrow points from the text ' $f(x)$  better be close to 1' to the first term. Another blue arrow points from the first term to the second term.

$$y_i \log(f(x)) + (1 - y_i) \log(1 - f(x))$$

If  $y_i = 0$



The diagram shows the term  $y_i \log(f(x)) + (1 - y_i) \log(1 - f(x))$  with a blue '0' to the left of the first term and a blue '1' below the second term. A blue arrow points from the text ' $f(x)$  better be close to 0' to the second term. Another blue arrow points from the first term to the second term.

$$y_i \log(f(x)) + (1 - y_i) \log(1 - f(x))$$

$f(x)$  better be close to 0

# How Do We Optimize This?

- We have an objective  $L(y, f(x))$  where
$$f(x) = \sigma(w_1x_1 + \cdots + w_nx_n + b)$$
- We can compute derivative of this function with respect to model parameters:  $L'(y, f(x))_{w,b}$
- Computed derivatives called *gradient*

# Computing Derivatives

- Sigmoid function:  $\sigma(x) = \frac{1}{1+e^{-x}}$   
 $\sigma'(x) = \sigma(x)(1 - \sigma(x))x'$

Thus


$$\begin{aligned}\sigma'_{w_1}(\dots) &= \sigma(\dots)(1 - \sigma(\dots))x_1 \\ \sigma'_b(\dots) &= \sigma(\dots)(1 - \sigma(\dots))\end{aligned}$$

- Logarithm function:  $\log(x)$   
 $\log'(x) = \frac{1}{x}x'$

# Optimization Procedure

- Okay, now we have derivatives with respect to parameters depending on our training data
- Let's make an arbitrary model (i.e. initialize *weights* randomly)
- Let's compute our loss  $L(y, f(x))$
- Then the derivative  $L'(y, f(x))_{w,b}$
- Then for each weight let's tweak it a little bit so our loss function becomes slightly better

$$w_i := w_i - \alpha L'(y, f(x))_{w_i}$$


 *Learning rate* – how strong should  
be an optimization step

# Optimization Procedure

1. Compute loss
  2. Compute the gradient for this loss w.r.t. the model parameters
  3. Update the model a little bit towards gradients
  4. Go to 1
- This is called *gradient descent*




# Gradient Descent Intuition




# Learning Rate

- *Learning rate* is a parameter you need to adjust wisely



Too small: converge  
very slowly



Too big: overshoot and  
even diverge

# Logistic Regression Recap

- Model output (0.5 value is a decision boundary)

$$f(x) = \sigma(w_1x_1 + w_2x_2 + \cdots + w_nx_n + b)$$

- Loss function (how bad things are)

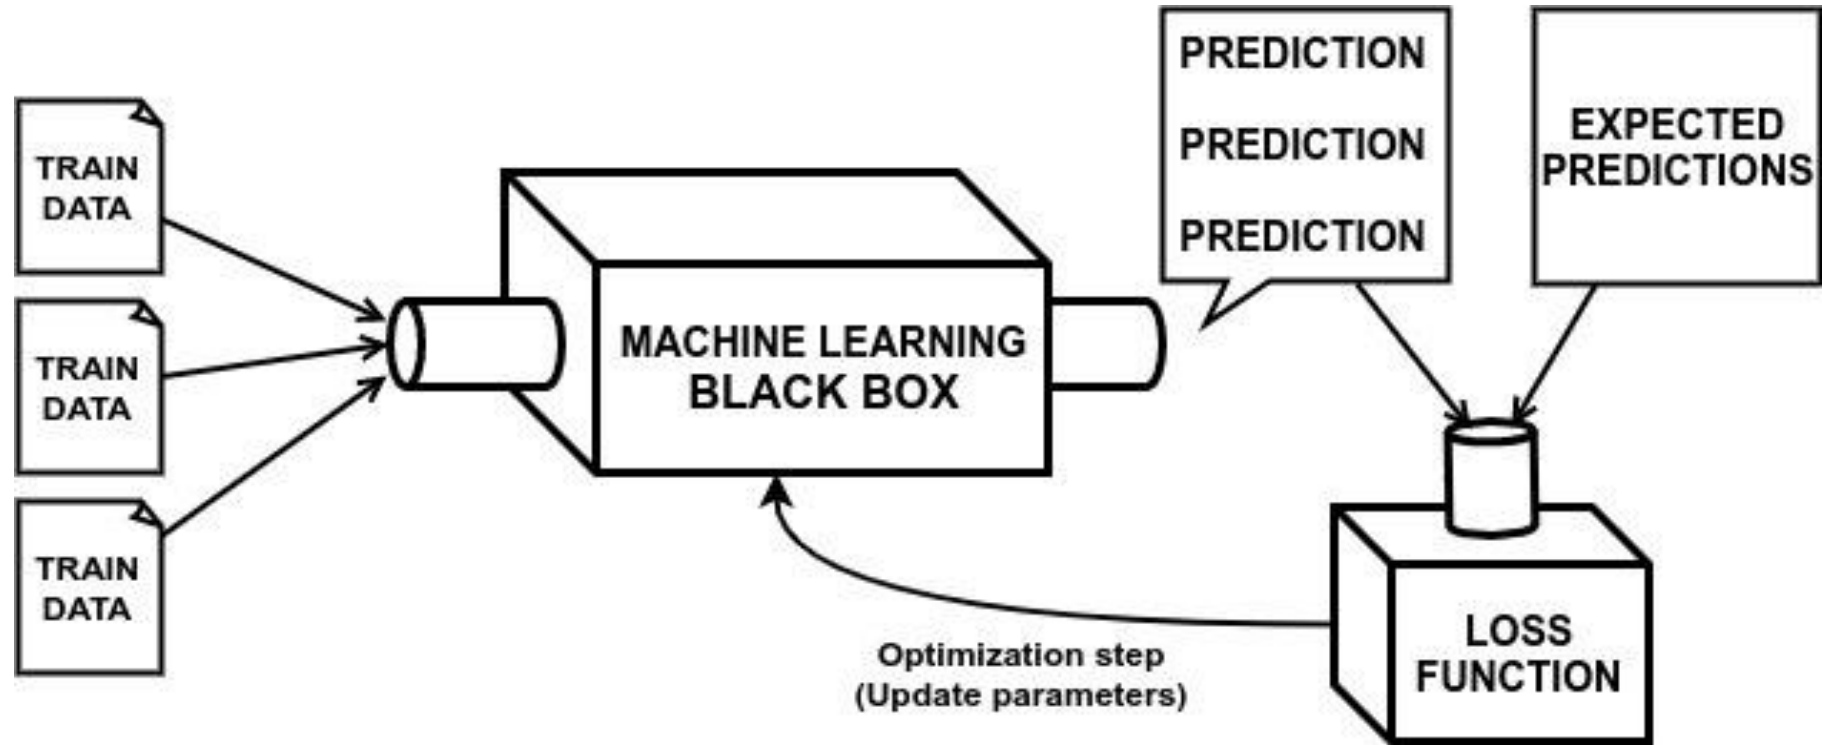
$$L(y, f(x)) = -\frac{1}{n} \sum_i^n [y_i \log(f(x)) + (1 - y_i) \log(1 - f(x))]$$

- We update parameters computing gradient and slightly changing the model

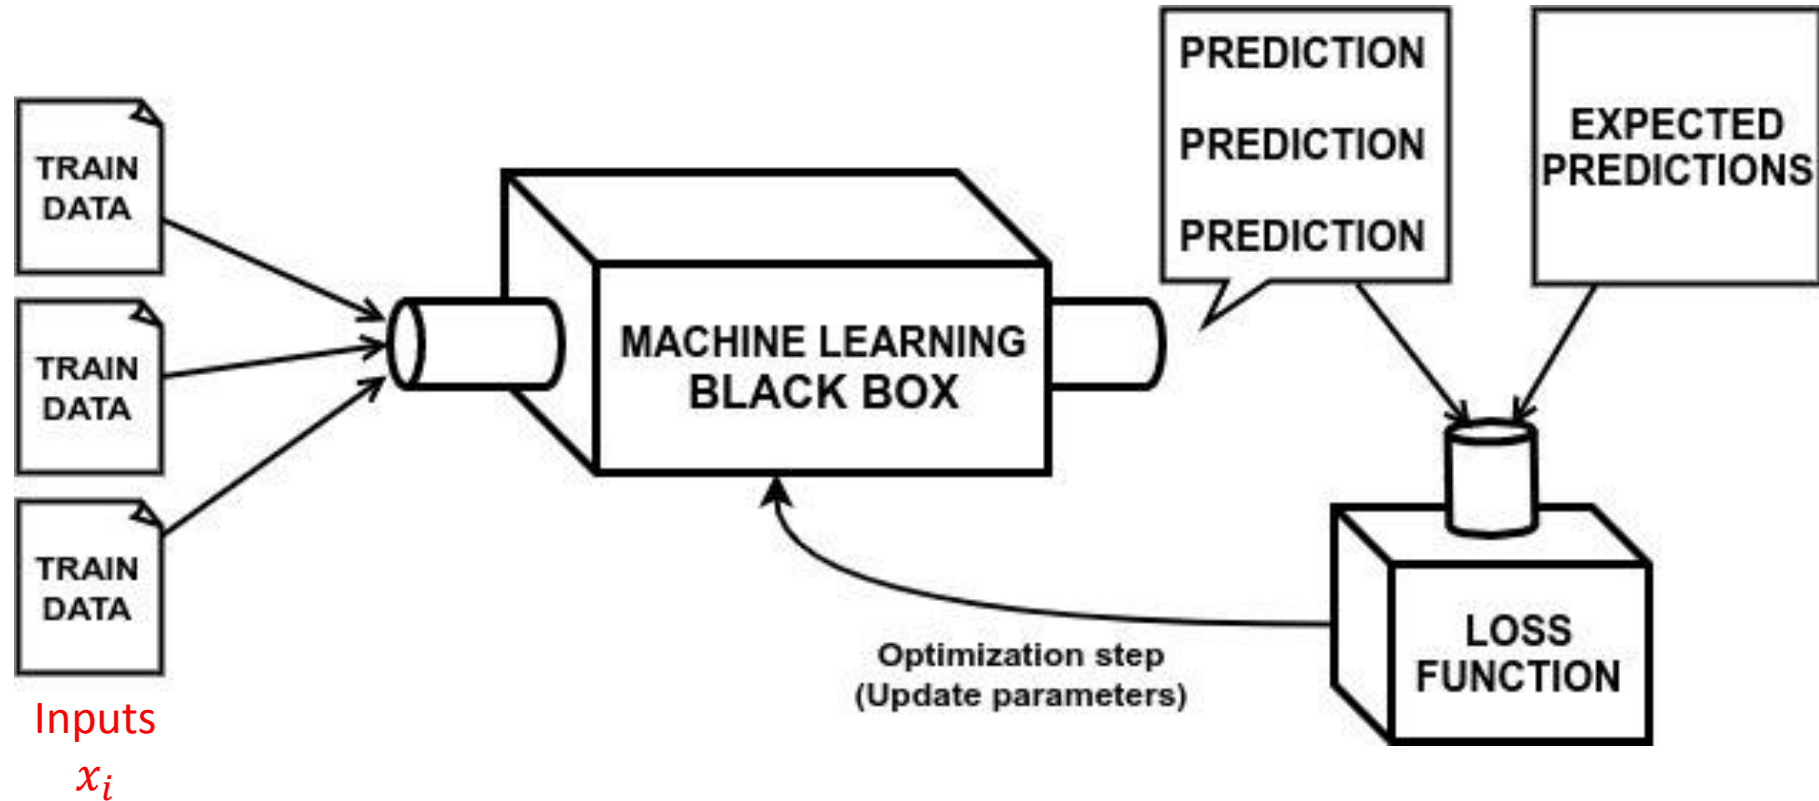
$$w_i := w_i - \alpha L'(y, f(x))_{w_i}$$

- Hoping the model will optimize itself in the best possible way

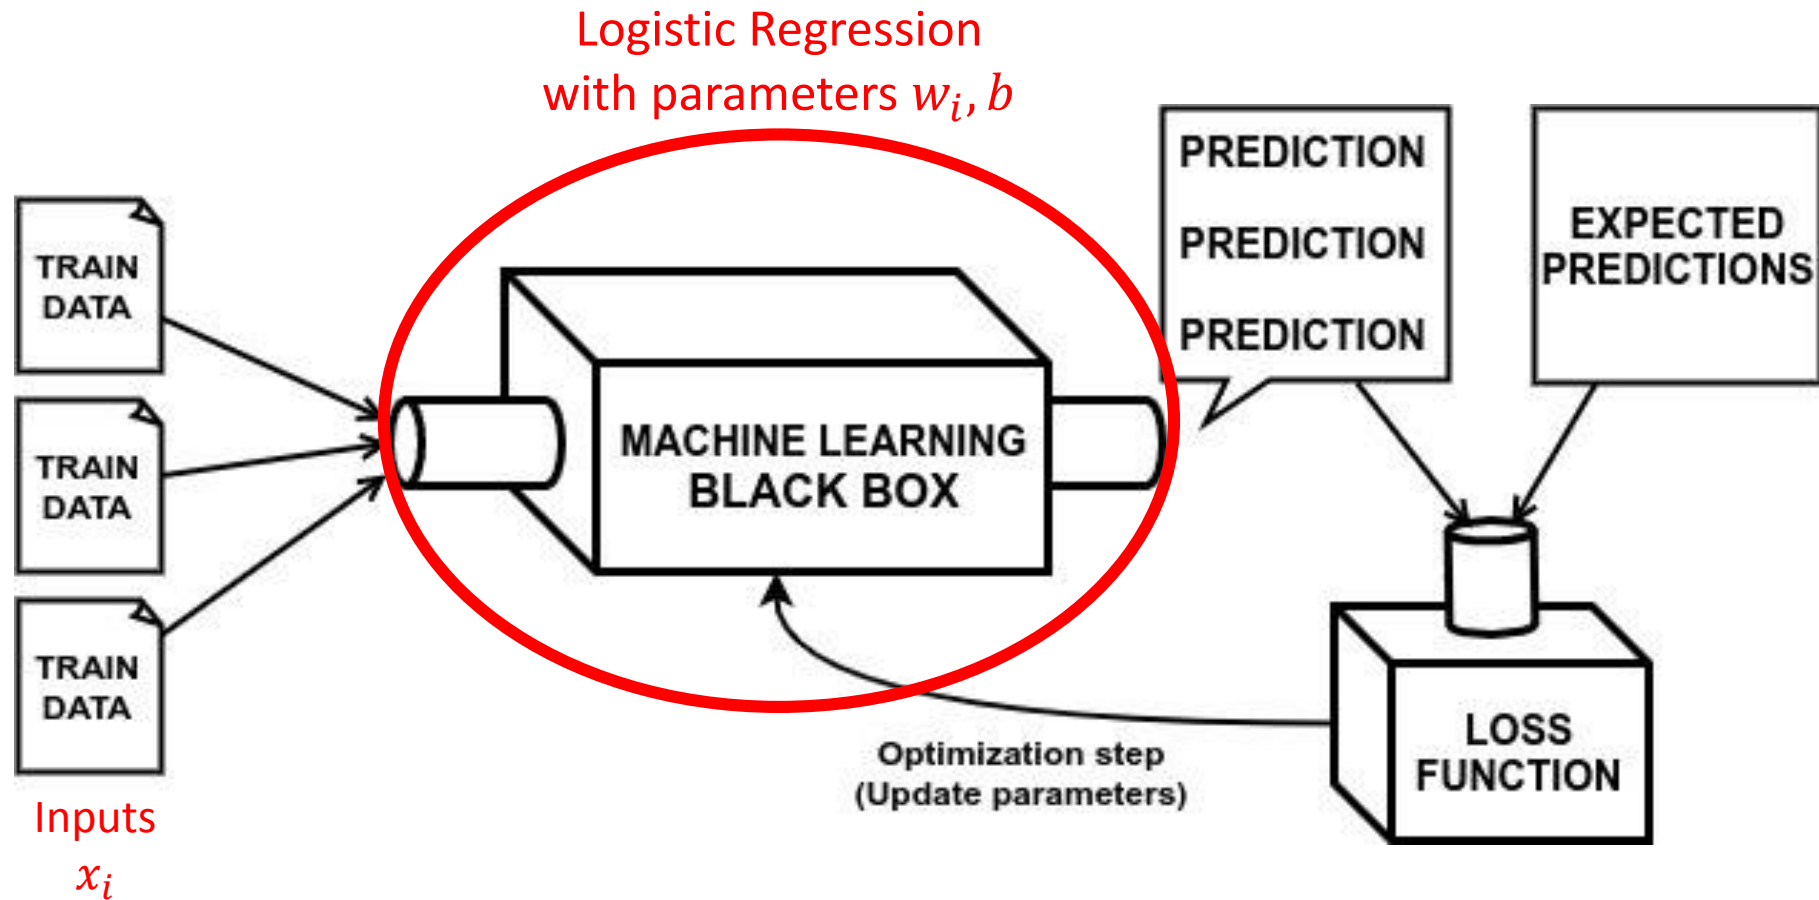
# Our Scheme



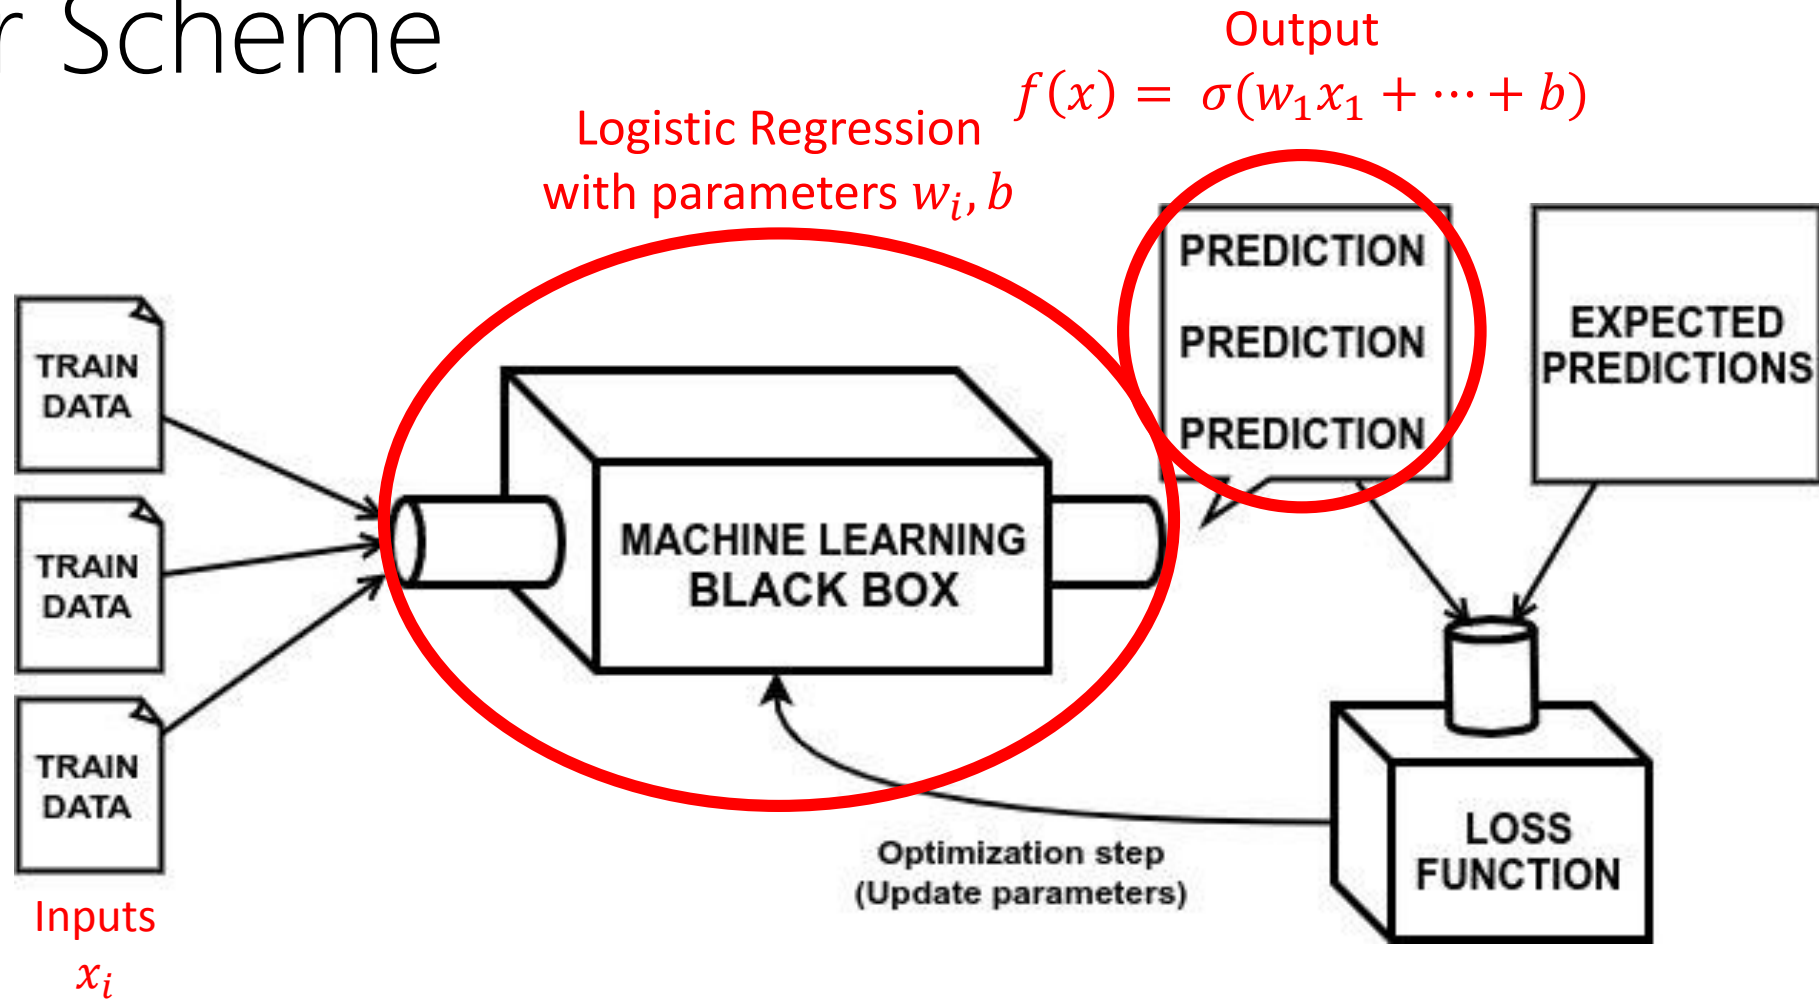
# Our Scheme



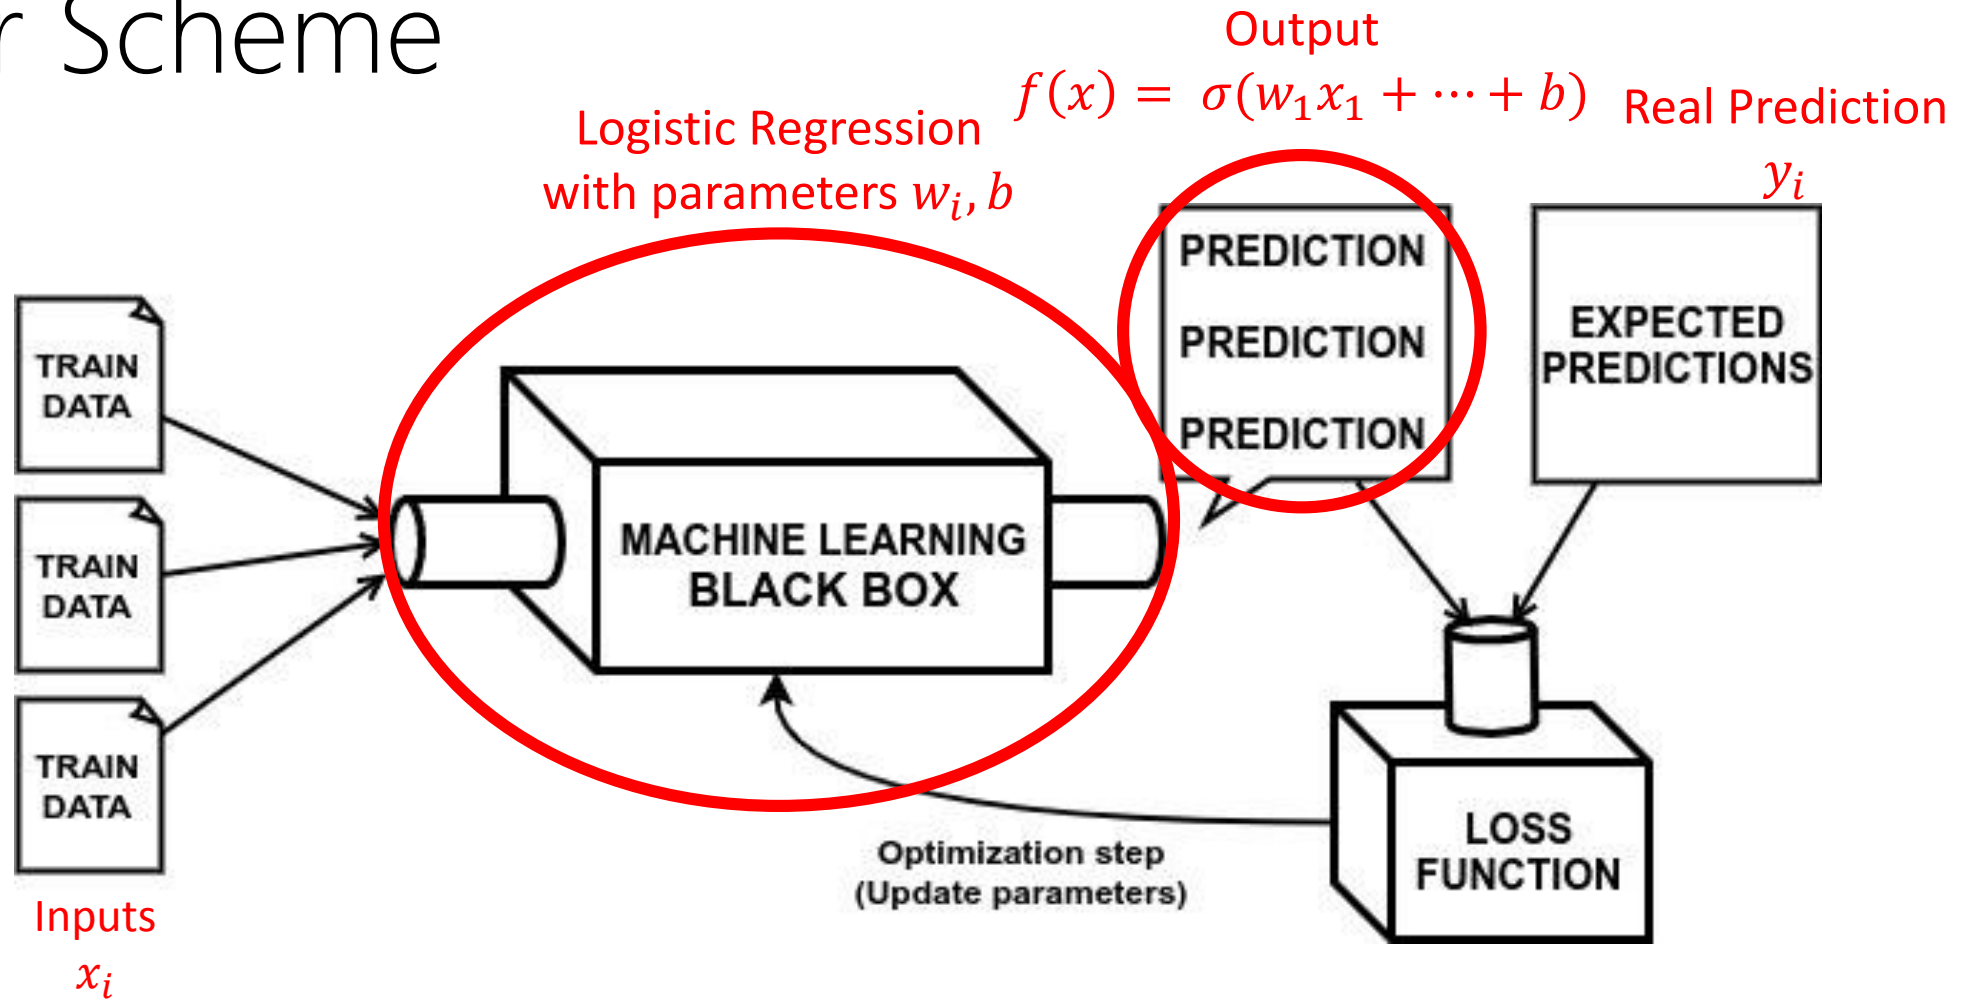
# Our Scheme



# Our Scheme

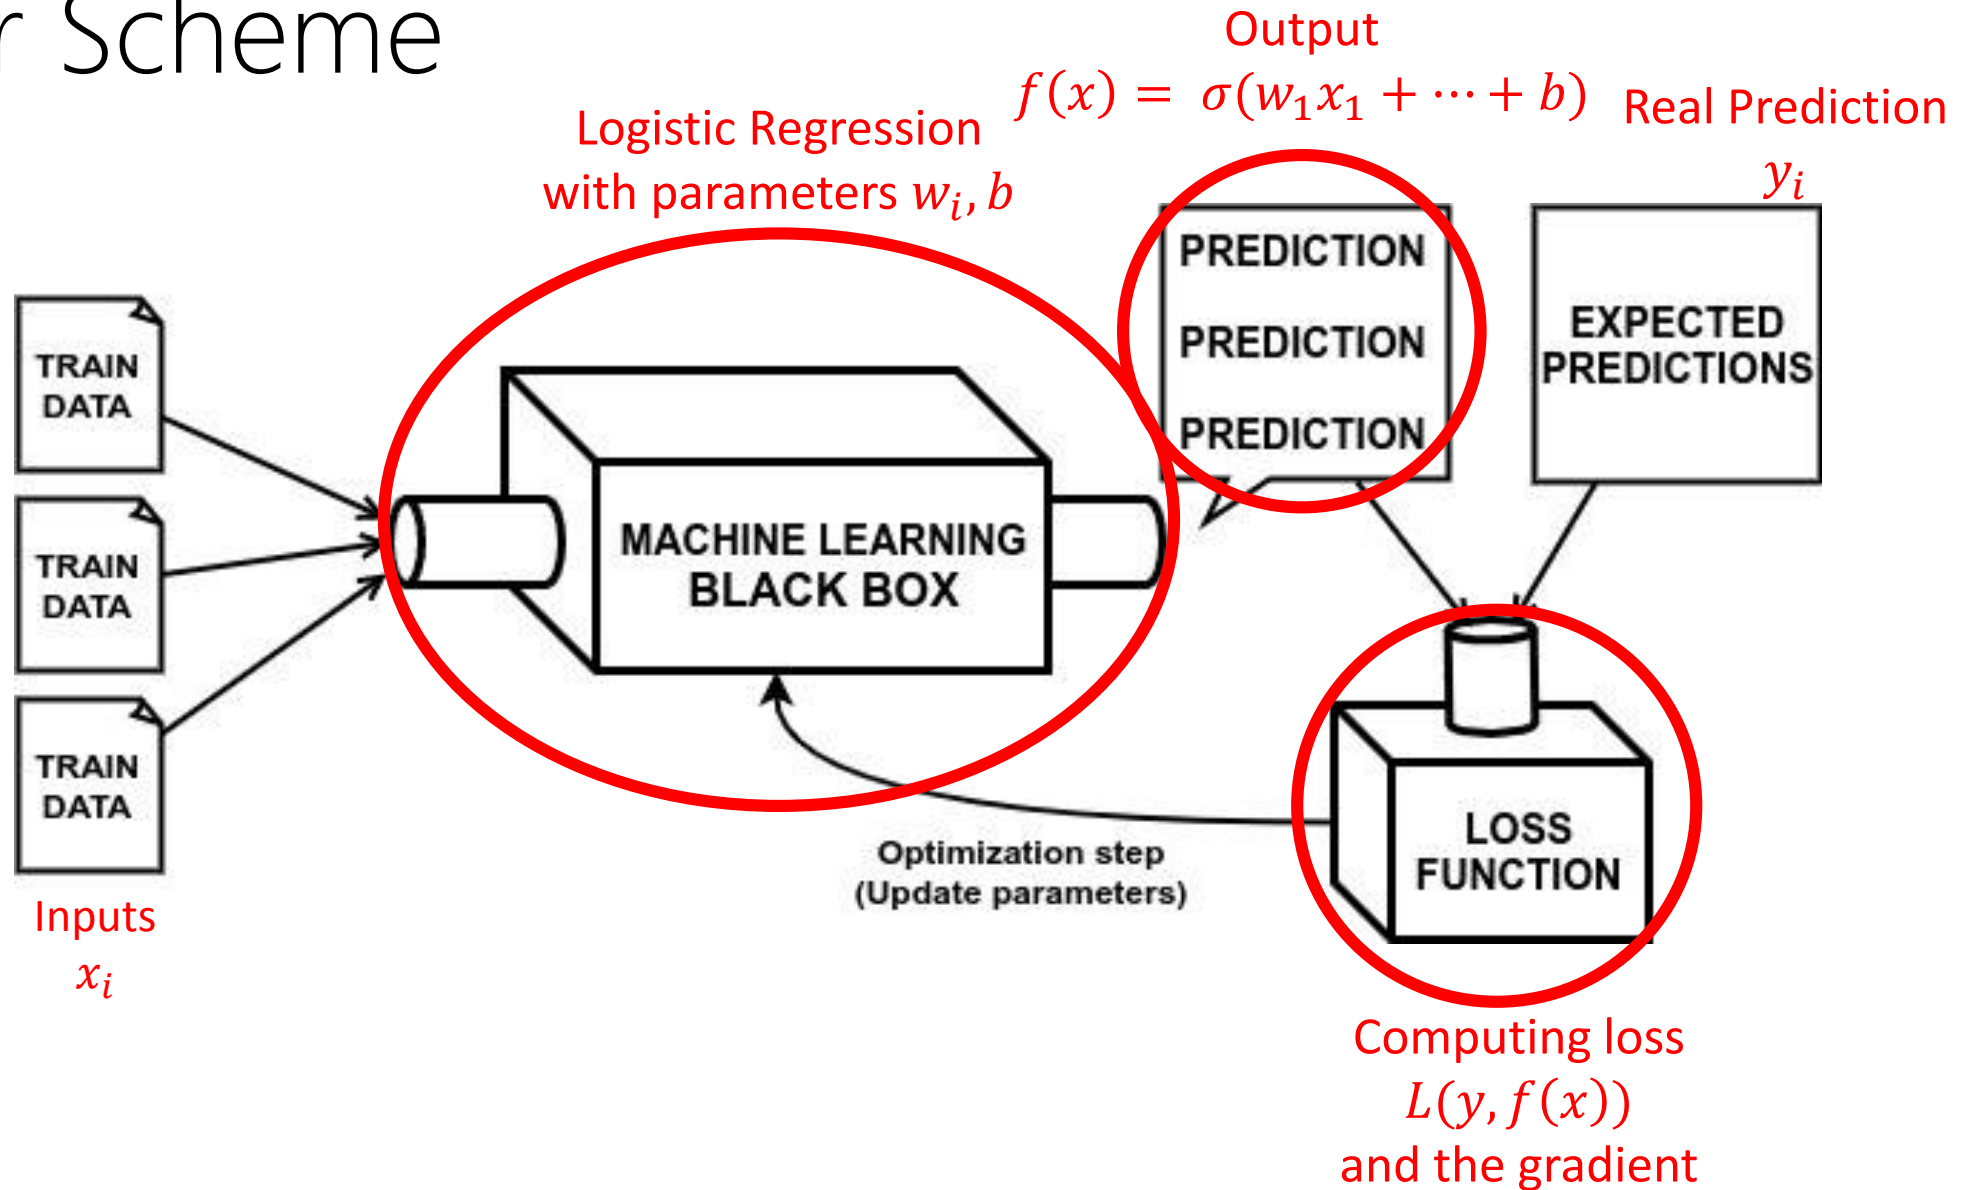


# Our Scheme

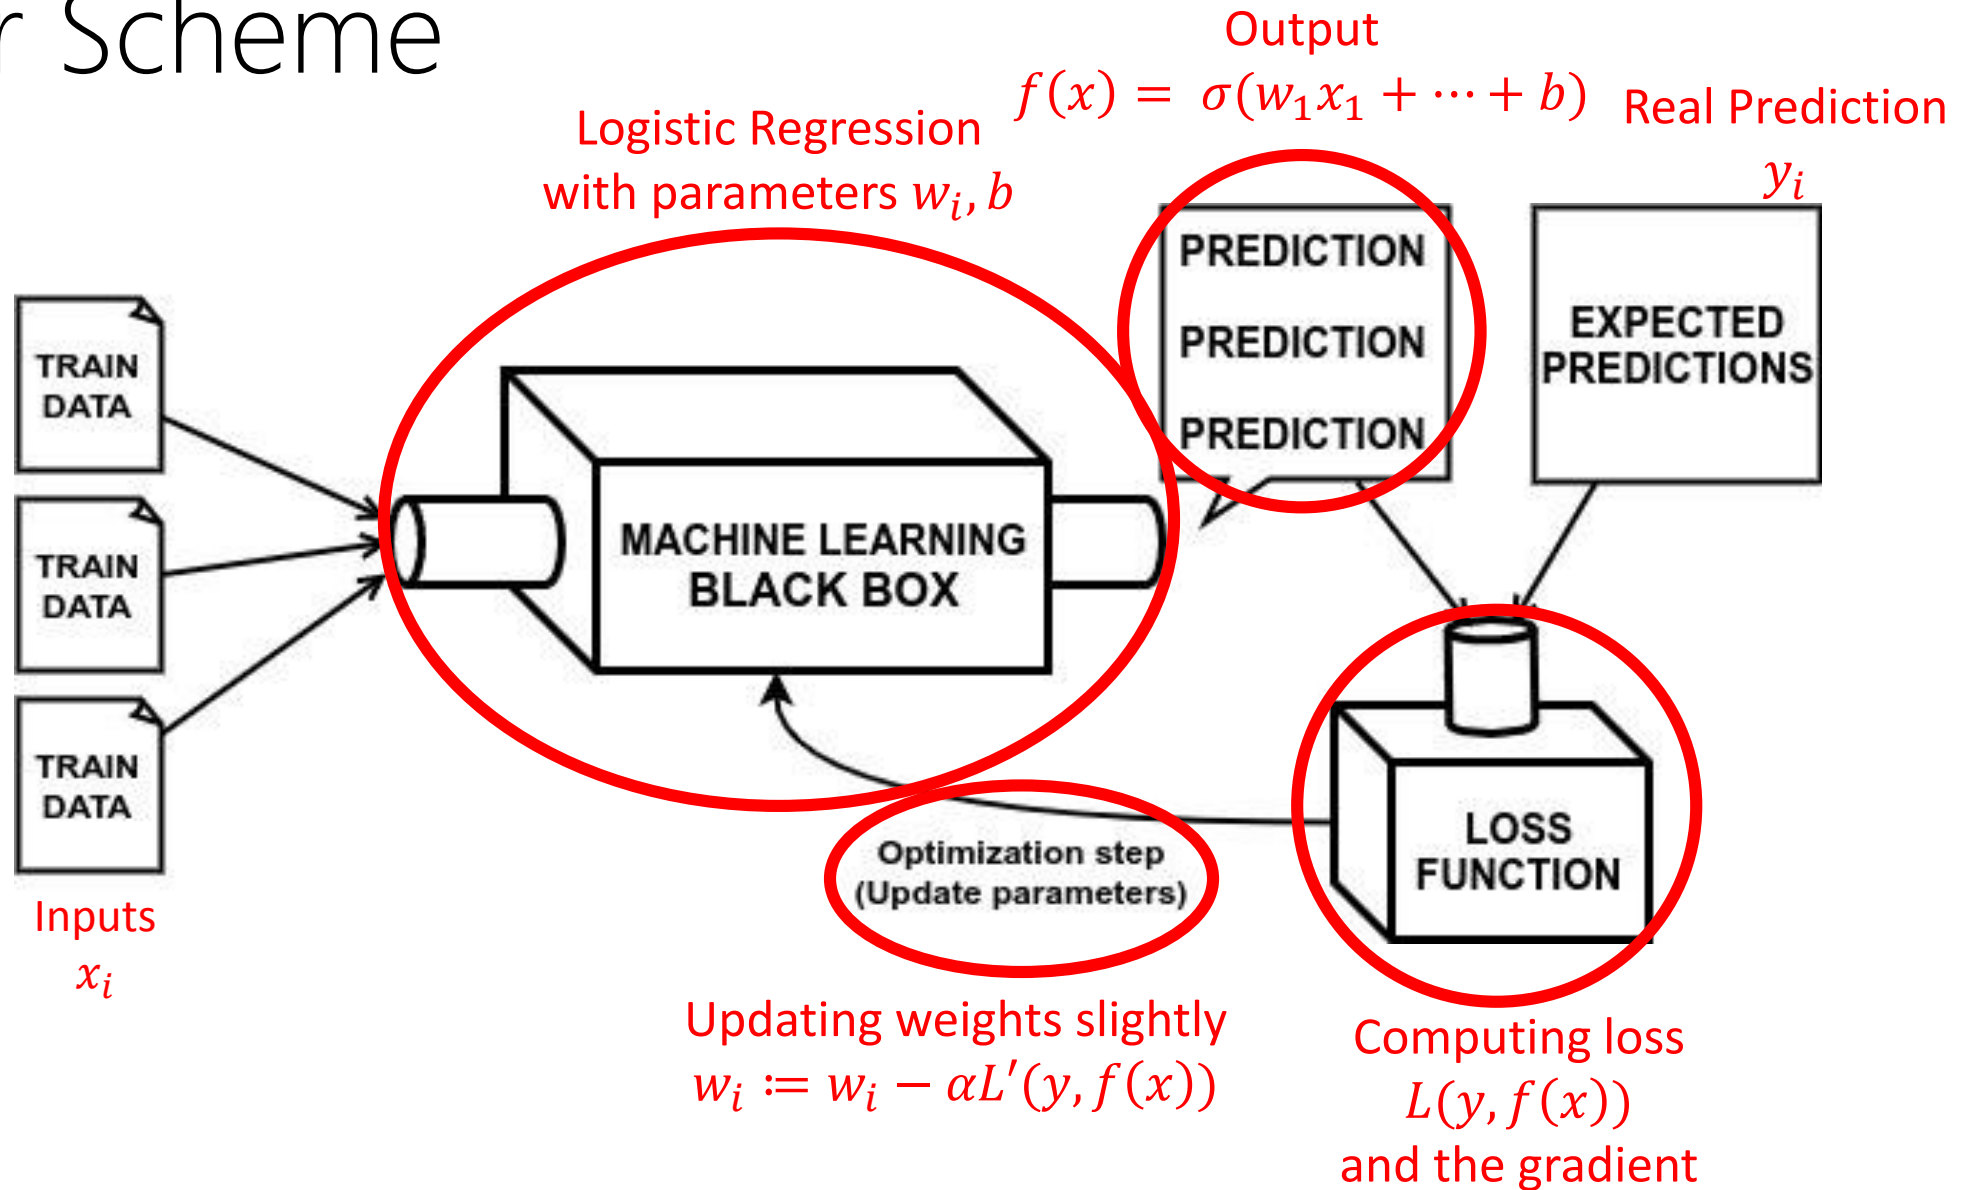





# Our Scheme




# Our Scheme




# Logistic Regression Scheme



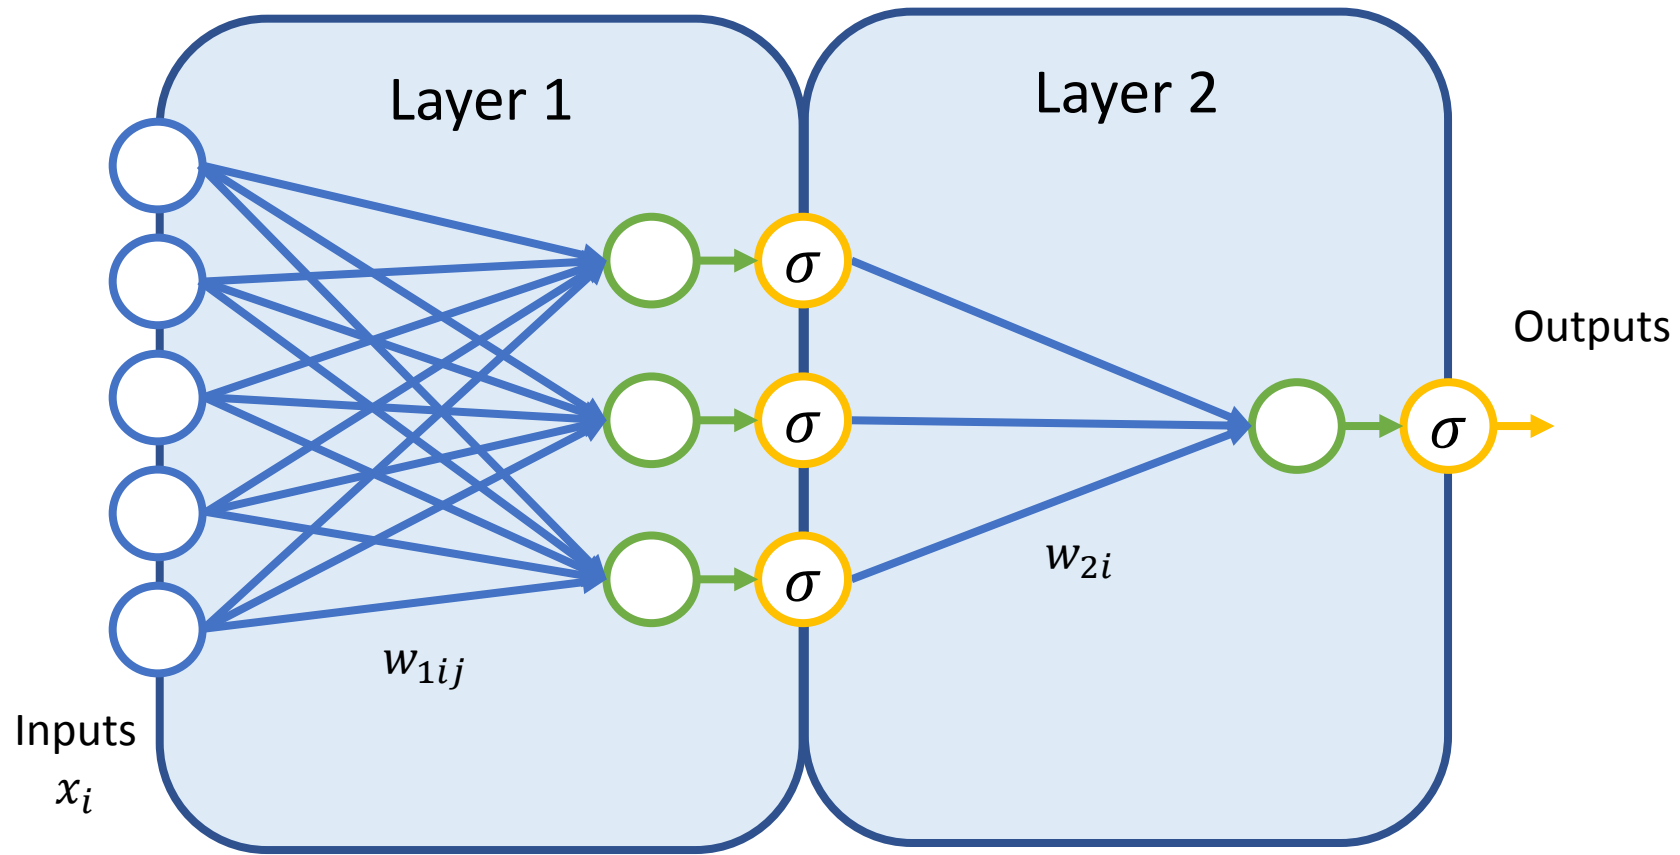
# Logistic Regression as a 1-layer NN




# Multi-output Layer



# Layers as Blocks



# Layers as Blocks



# Neural Networks

- You can stack many layers with various size of outputs together to achieve more complex models
- Don't forget about *non-linearity* at the end of the layers (this is crucial)
- The deeper network – the more complex features *layers* can learn



# Lecture Recap

- ML is really powerful nowadays and grows rapidly
- Splitting data into *train* and *validation* subsets to control overfitting
- Logistic Regression learns to separate data with a line (or a plane)
- Optimization of Logistic Regression with *gradient descent*
- Picking *learning rate* carefully
- Logistic Regression is a 1-layer Neural Network
- You can stack multiple layers to get deeper and more effective NNs

This is it for the first lecture