


# Deep Neural Networks and Where to Find Them

Lecture 2

Artem Korenev, Nikita Gryaznov


Recap

# Recap

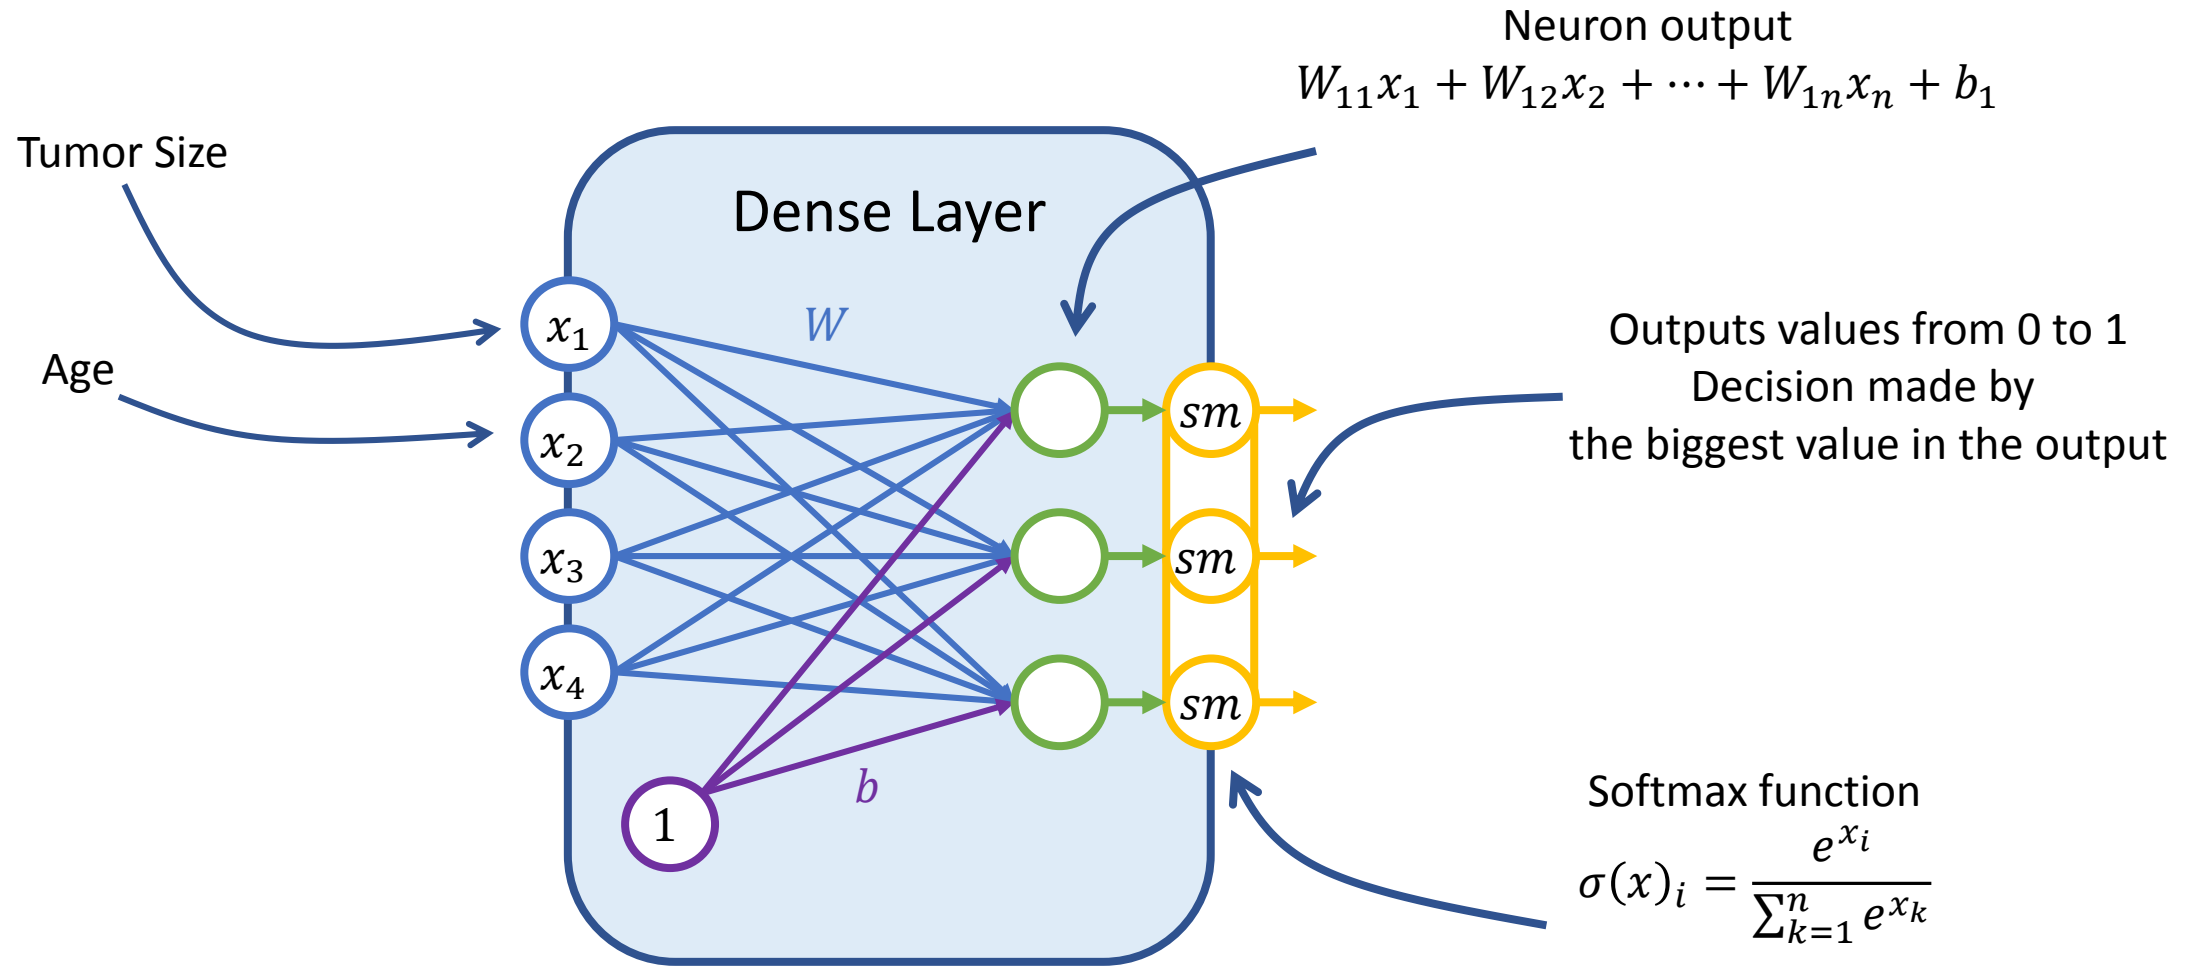


- $x_1$  Layer input
- Neuron Output  
( $W_{i1}x_1 + \dots + W_{i2}x_2 + b_i$ )
- $f$  Activation function  
(Non-linearity)  
(Neuron activation)
- Layer weight  
(trainable parameter)
- Layer bias  
(trainable parameter)


# Binary Logistic Regression




# Multiclass Logistic Regression




## 2 Layer NN with one output



# 3 Layer NN with two outputs



# Logistic Regression Optimization






# Logistic Regression Optimization

Compute gradient for *weights*

“Given our data, how should we change  $w_1$  slightly so Loss function value is a little bit better?”




# Logistic Regression Optimization

Compute gradient for *weights*

“Given our data, how should we change  $w_1$  slightly so Loss function value is a little bit better?”

Then we update all *weights* slightly towards the gradient hoping that loss function value will improve

$$w_i := w_i - \alpha L'(y, \sigma(\dots))_{w_i}$$



# Logistic Regression Optimization

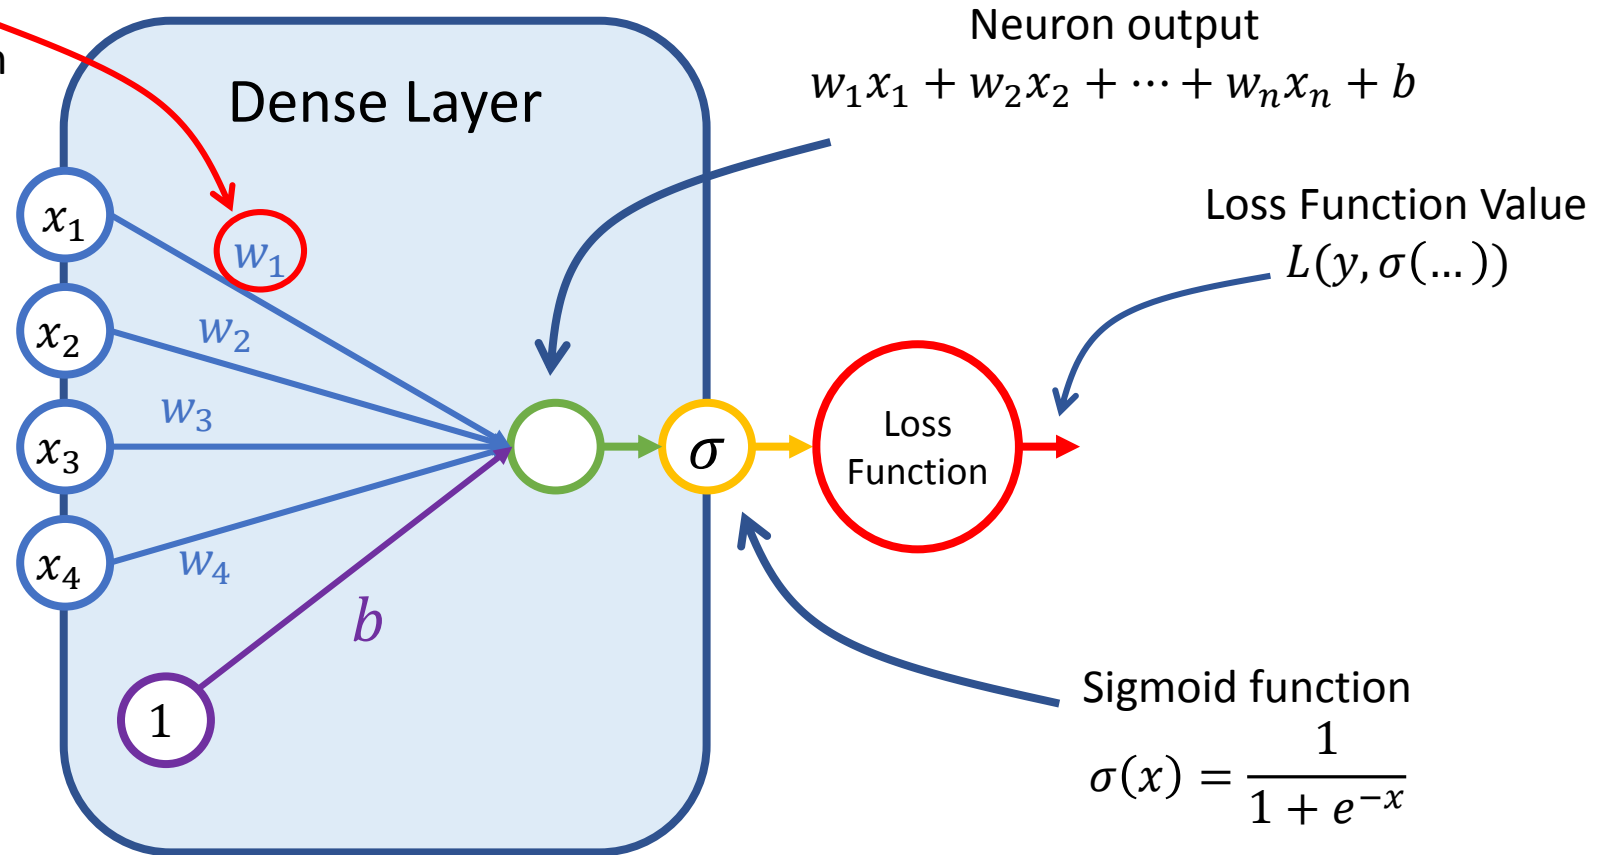
Compute gradient for *weights*

“Given our data, how should we change  $w_1$  slightly so Loss function value is a little bit better?”

Then we update all *weights* slightly towards the gradient hoping that loss function value will improve

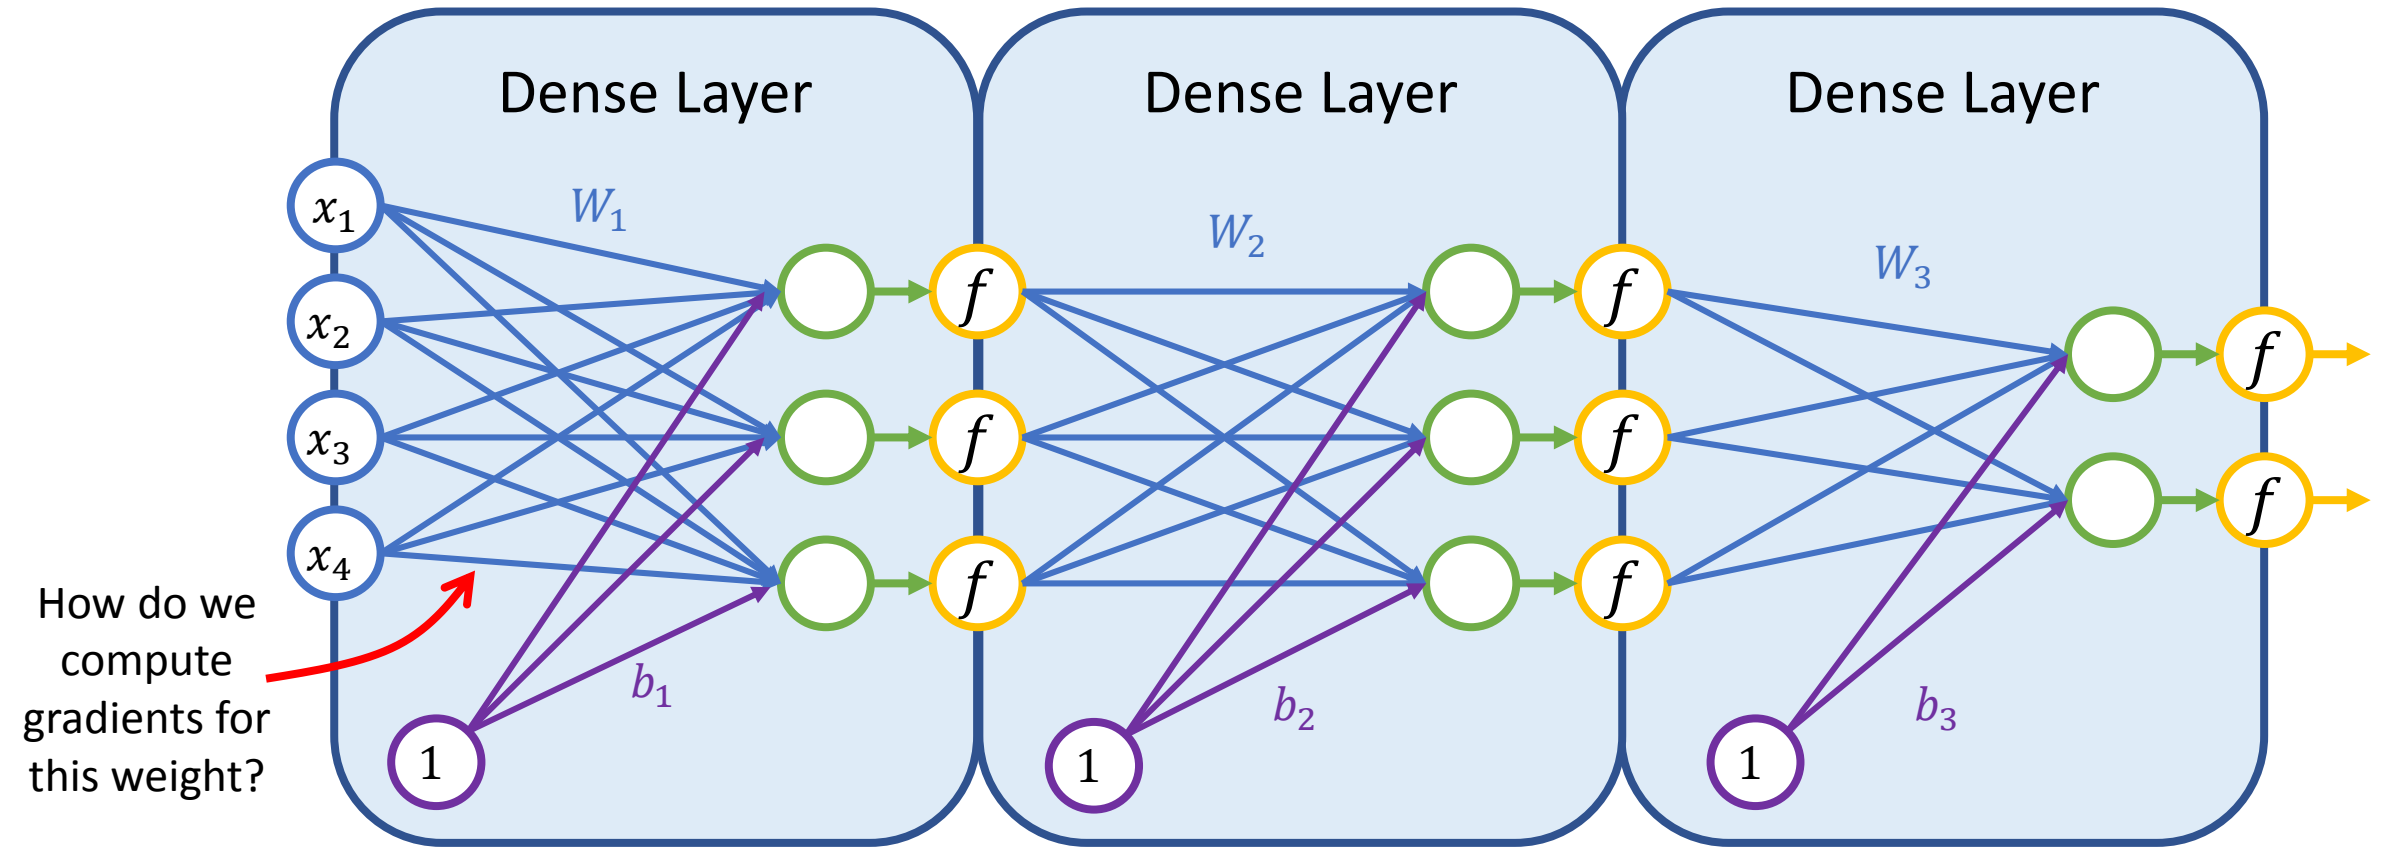
$$w_i := w_i - \alpha L'(y, \sigma(\dots))_{w_i}$$

Repeat calculating loss, computing gradients and updating the *weights* (*gradient descent*)



NN Optimization. Chain Rule

# Optimizing Complex Networks



# How to update weights?

- We need to compute gradients for deeper layers
- We can do it viewing our neural network as a *computational graph* and using *the chain rule*
- Computing gradients and optimizing the neural network called *backpropagation*

# Chain rule

$$L = f(g(x))$$

$$\frac{dL}{dx} = \frac{df}{dg} \frac{dg}{dx}$$

$$L = \log(x^2)$$


$$f(x) = \log(x) \quad g(x) = x^2$$

$$\frac{dL}{dx} = \frac{1}{x^2} \times 2x = \frac{2}{x}$$

# Chain rule

Assume that we did forward pass and evaluated values at each node


$$L = z(2x + y)$$





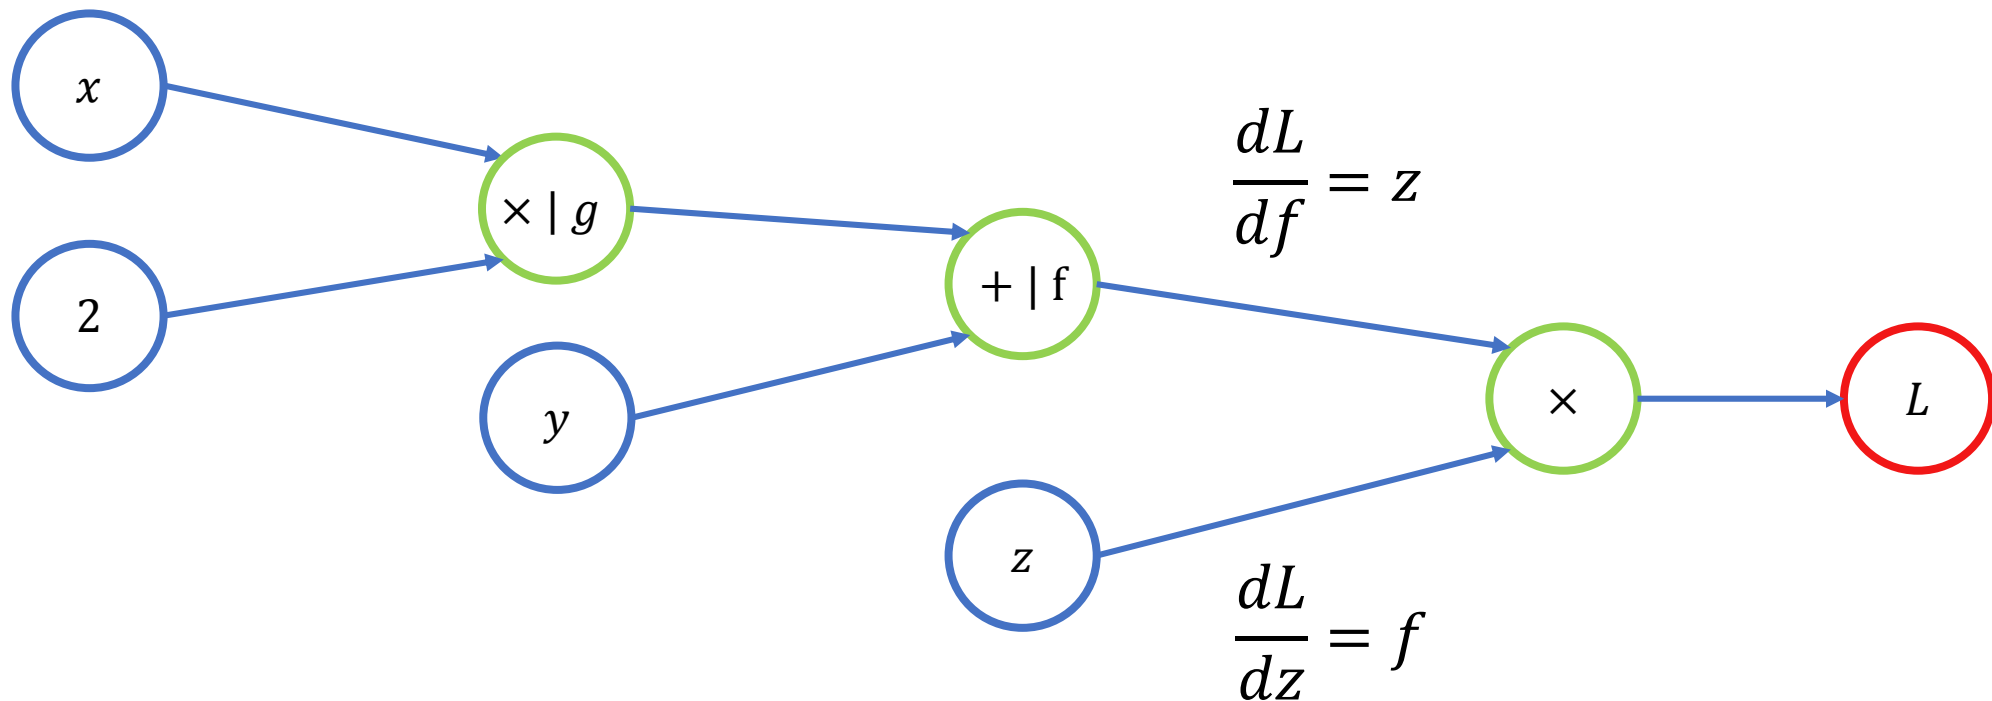
# Chain rule

$$\begin{aligned}g(x) &= 2x \\f(g, y) &= g + y \\L(z, f) &= zf\end{aligned}$$




# Chain rule

$$\begin{aligned}g(x) &= 2x \\f(g, y) &= g + y \\L(z, f) &= zf\end{aligned}$$



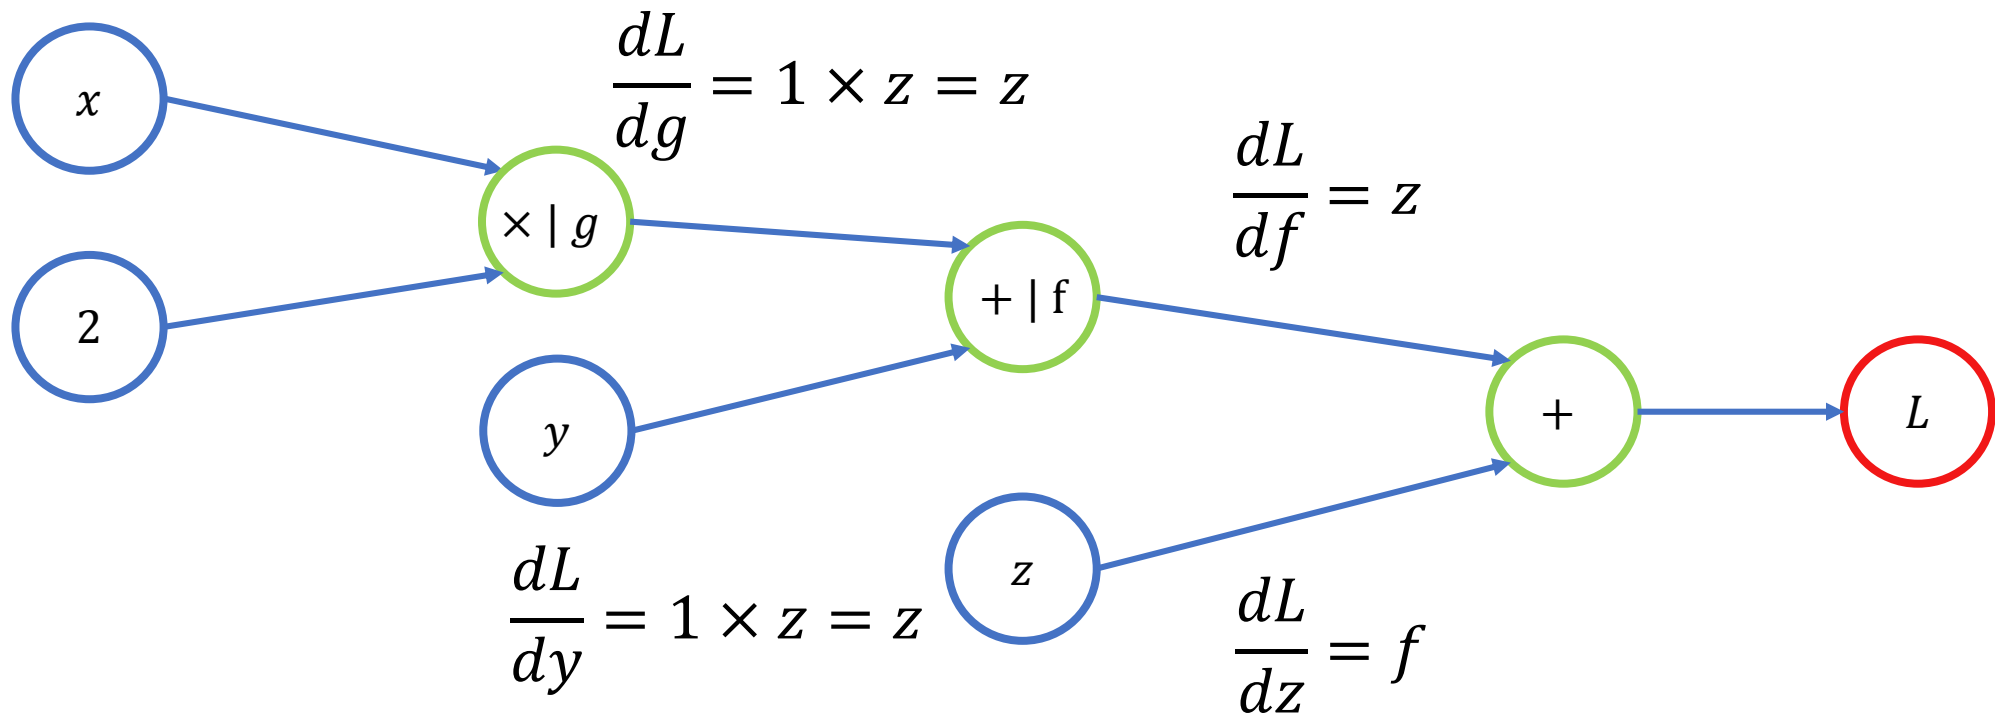
# Chain rule

$$\begin{aligned}g &= 2x \\ f &= g + y \\ L &= zf\end{aligned}$$

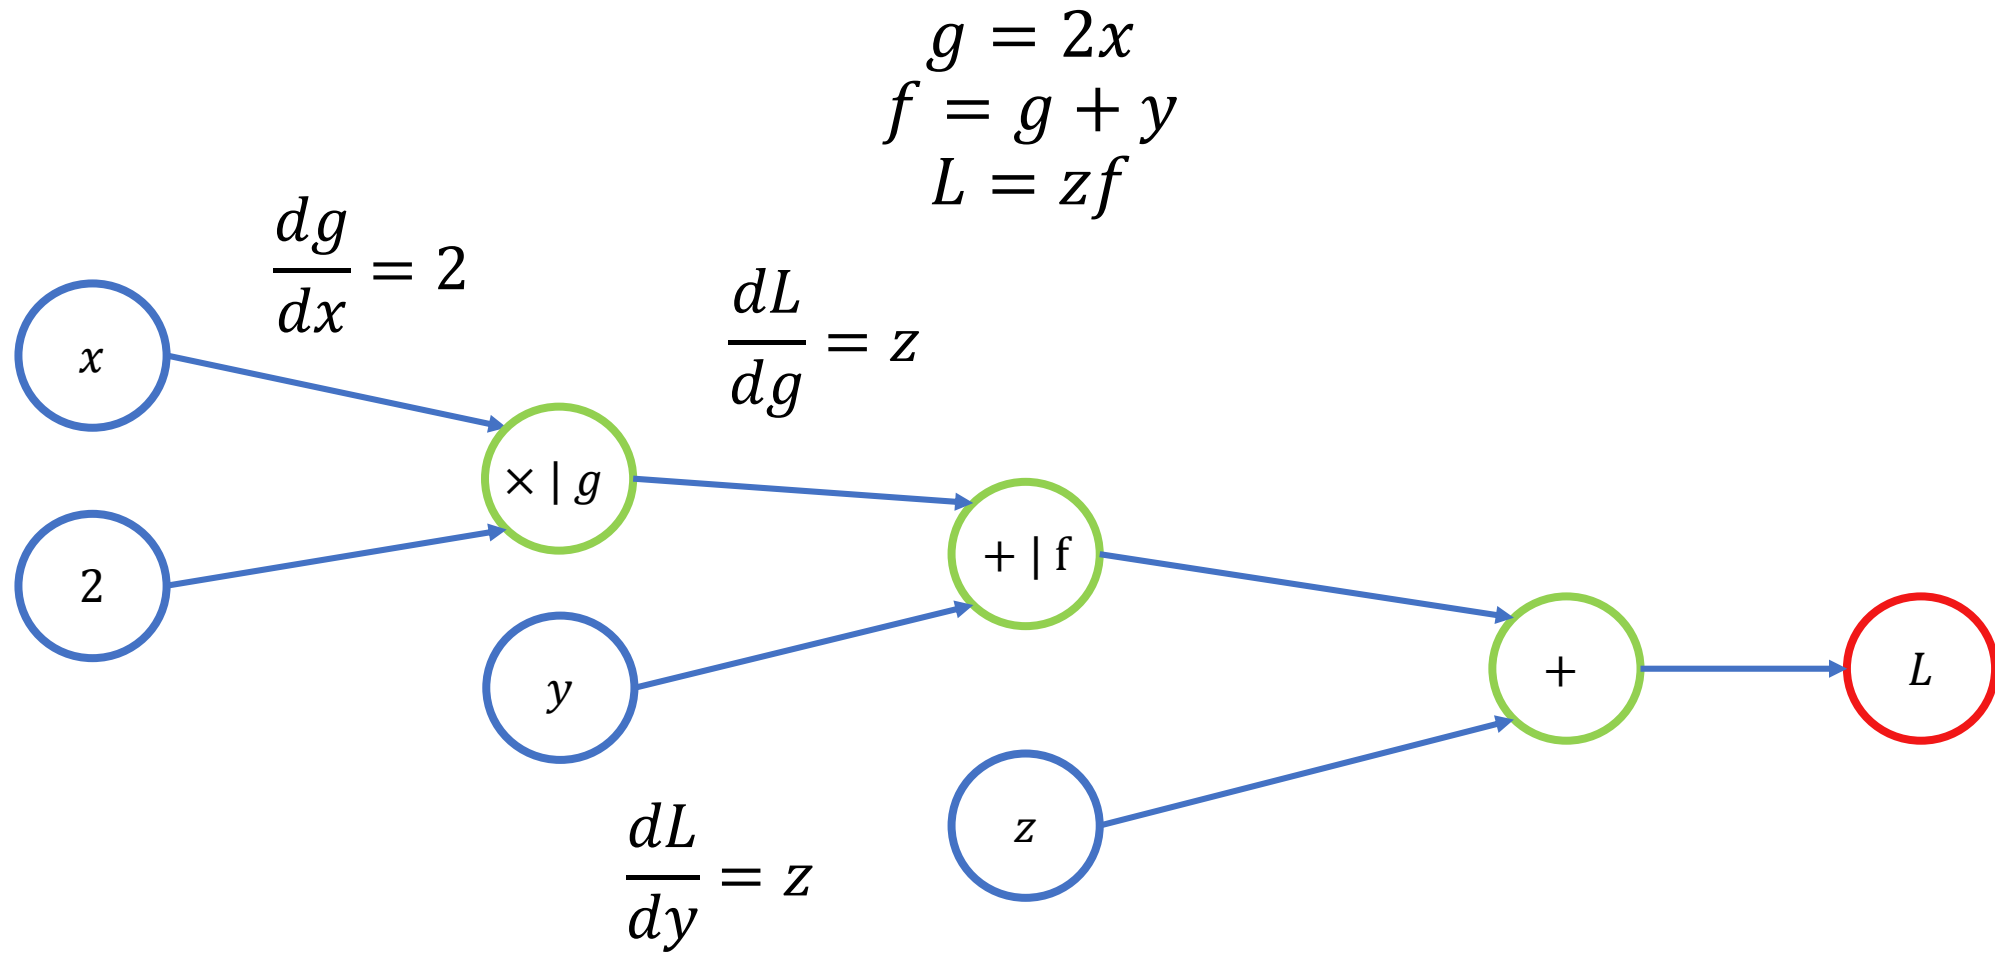


# Chain rule

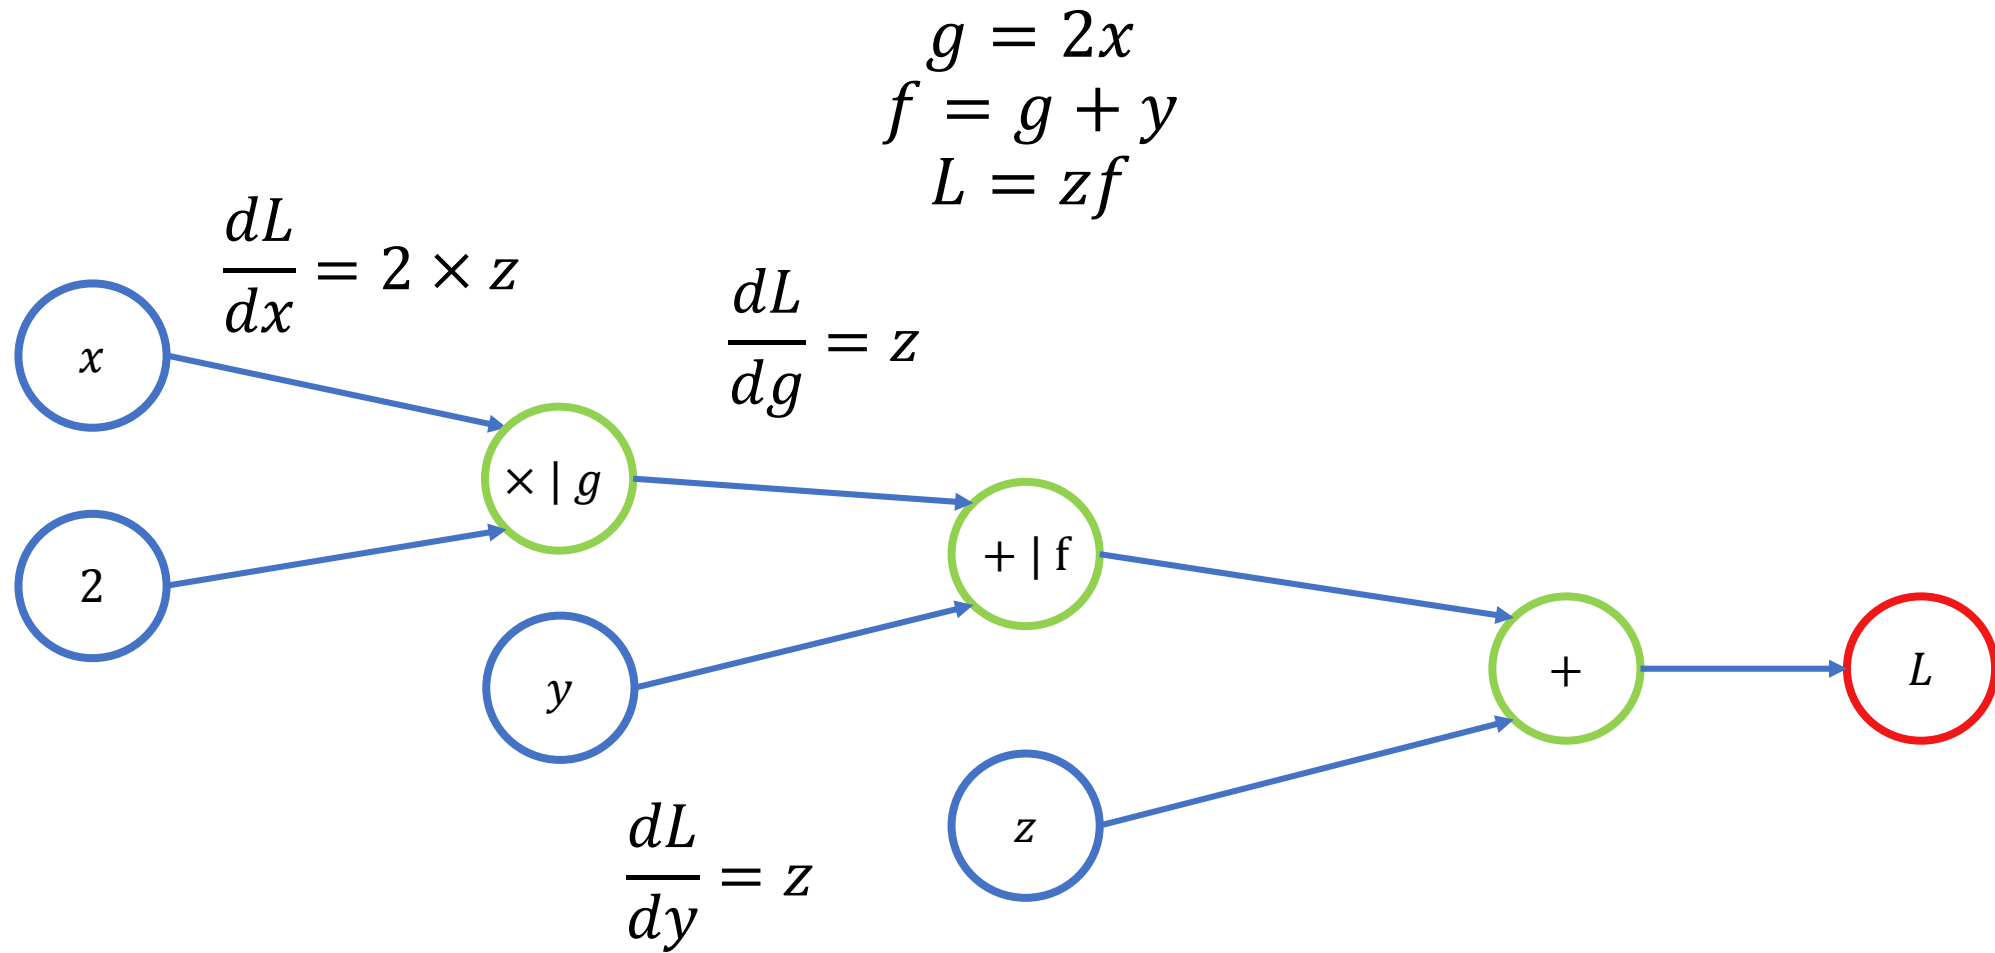
$$\begin{aligned}g &= 2x \\ f &= g + y \\ L &= zf\end{aligned}$$



# Chain rule




# Chain rule

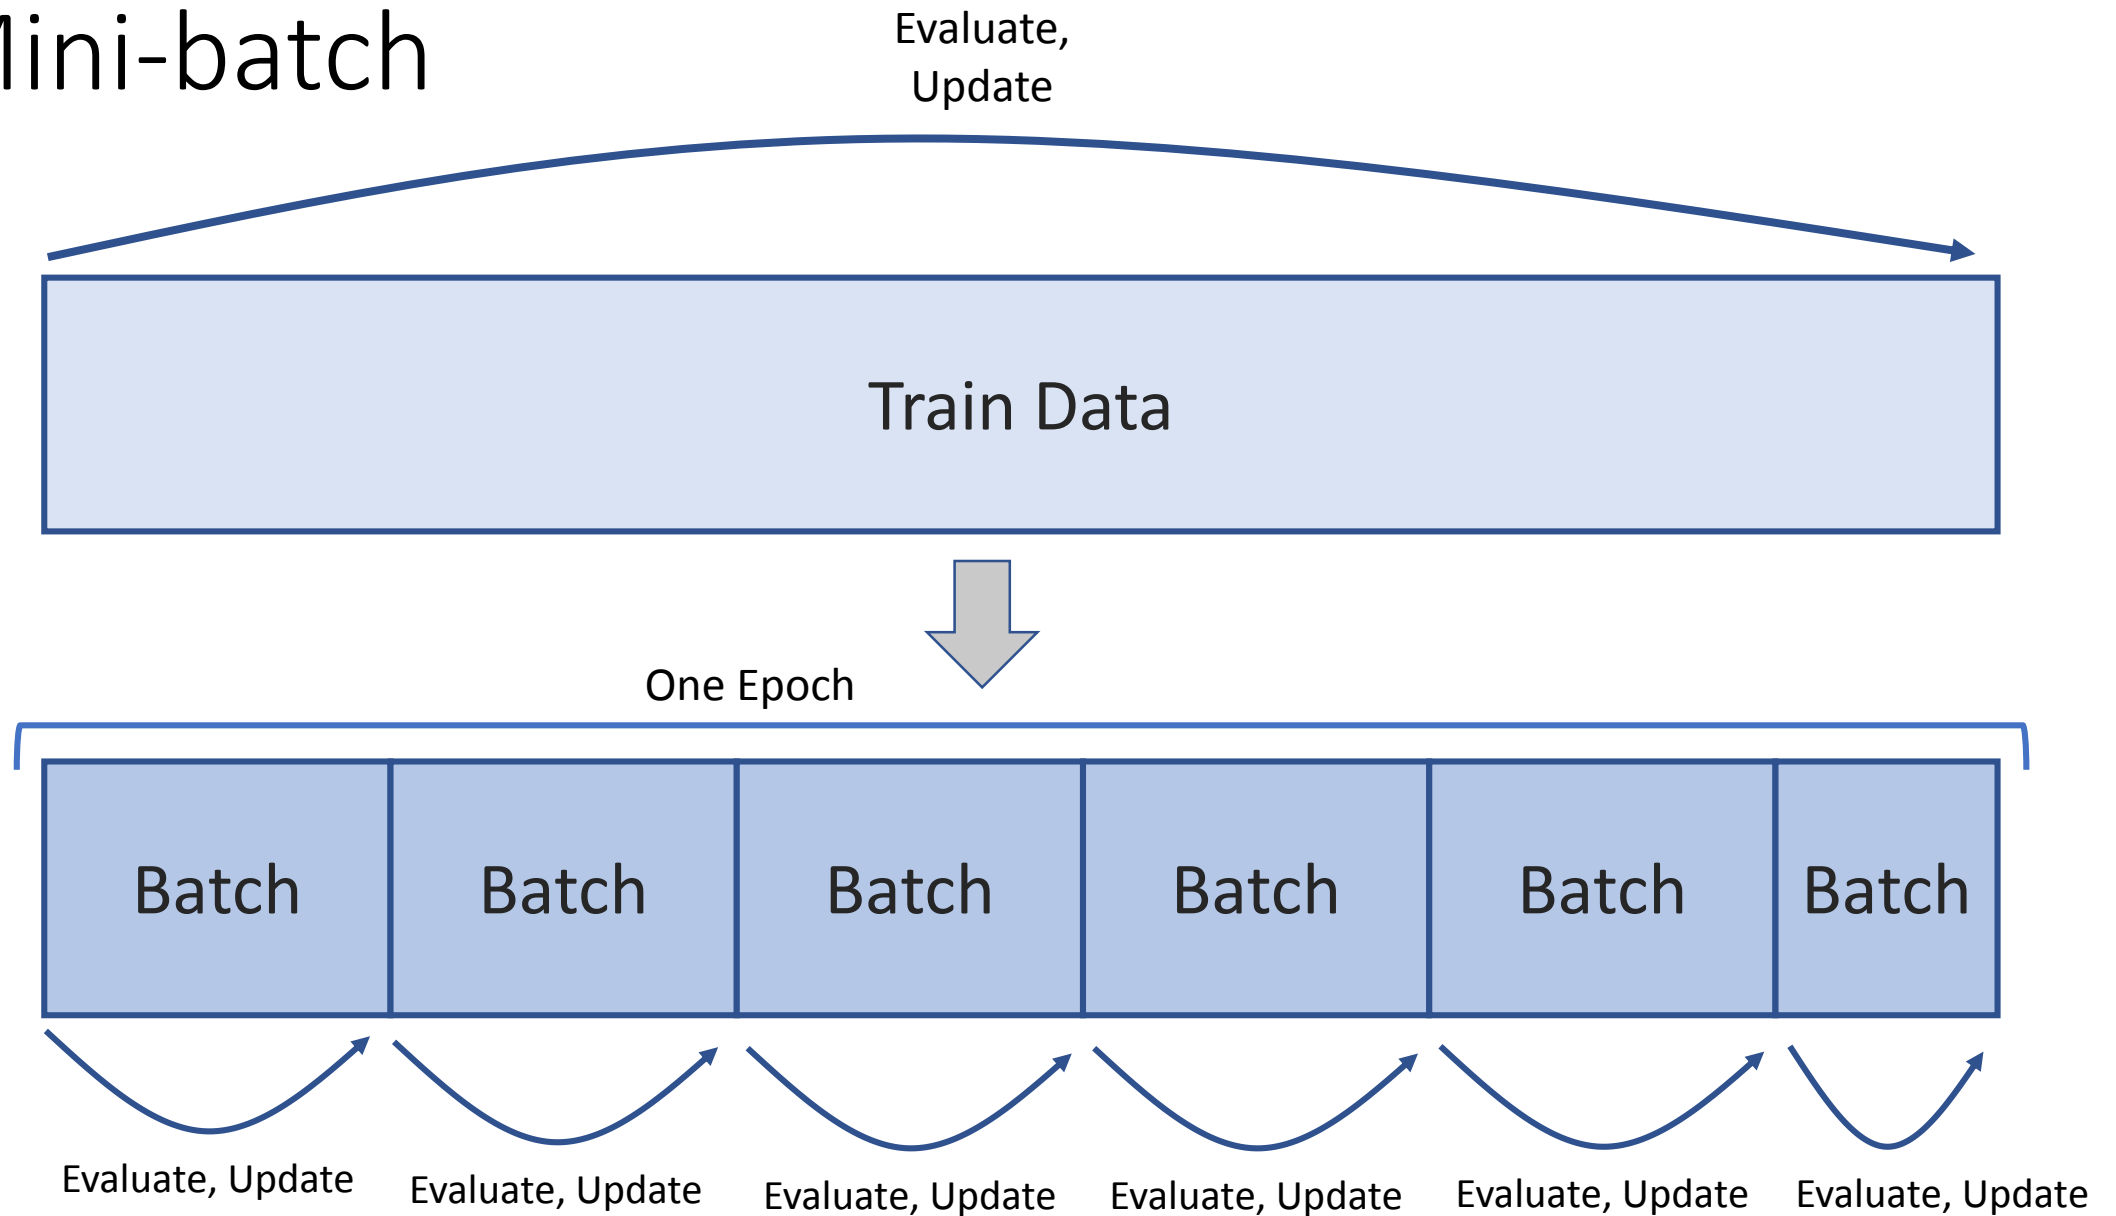


# Backpropagation

At each step we only need to store gradient of the next layer and values of its inputs



# Mini-batch





# NN Matrix Representation

# Linear Algebra Recap

- Vector – a list of numbers

$$b = \begin{pmatrix} 1 \\ 7 \\ 5 \end{pmatrix}$$

- Matrix – a 2-dimensional list of numbers

$$A = \begin{pmatrix} 6 & 2 & 3 \\ 1 & 5 & 7 \end{pmatrix}$$

- Matrix  $A$  have sizes (shape)  $A.shape = (2, 3)$
- Vector  $b$  is a matrix of shape  $b.shape = (1, 3)$
- Tensor is just another name for multi-dimensional matrix

# Vector and Matrix Operations

- Transposition (flipping)

$$b = \begin{pmatrix} 1 \\ 7 \\ 5 \end{pmatrix}, b^T = (1, 7, 5)$$

$$A = \begin{pmatrix} 6 & 2 & 3 \\ 1 & 5 & 7 \end{pmatrix}, A^T = \begin{pmatrix} 6 & 1 \\ 2 & 5 \\ 3 & 7 \end{pmatrix}$$

- Vector scalar multiplication

$$x = \begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix}, z = \begin{pmatrix} 4 \\ 5 \\ 6 \end{pmatrix}, x^T z = 1 * 4 + 2 * 5 + 3 * 6 = 32$$

# Vector and Matrix Operations

- Matrix-vector multiplication

$$A = \begin{pmatrix} 6 & 2 & 3 \\ 1 & 5 & 7 \end{pmatrix}, x = \begin{pmatrix} 1 \\ 7 \\ 5 \end{pmatrix}$$
$$Ax = \begin{pmatrix} 6 * 1 + 2 * 7 + 3 * 5 \\ 1 * 1 + 5 * 7 + 7 * 5 \end{pmatrix} = \begin{pmatrix} 35 \\ 71 \end{pmatrix}$$

- Matrix-matrix multiplication

$$A = \begin{pmatrix} 6 & 1 \\ 2 & 5 \\ 3 & 7 \end{pmatrix}, B = \begin{pmatrix} 0 & 2 \\ 1 & 3 \end{pmatrix}$$
$$AB = \begin{pmatrix} 0 * 6 + 1 * 1 & 6 * 2 + 1 * 3 \\ 0 * 2 + 1 * 5 & 2 * 2 + 5 * 3 \\ 3 * 0 + 7 * 1 & 2 * 3 + 7 * 3 \end{pmatrix} = \begin{pmatrix} 1 & 15 \\ 5 & 19 \\ 7 & 27 \end{pmatrix}$$

# Vector and Matrix Operations


- Matrix-matrix multiplication
- To multiply matrix  $A$  and  $B$ , last dimension of  $A$  and first dimension of  $B$  must match

$$\begin{aligned} A.shape &= (6, 7), & B.shape &= (7, 10), \\ (AB).shape &= (6, 10) \end{aligned}$$


# Input as a Matrix

- Our NN has input of size  $N$  (tumor size, age, ...)
- We pass  $B$  (batch size) data points at once
- We can represent our input batch as a matrix of shape  $(B, N)$
- Let's think of Dense layer matrix representation

# Dense Layer via Matrices




# Dense Layer via Matrices






# Dense Layer via Matrices



# Dense Layer via Matrices



# Dense Layer via Matrices

- Number of inputs:  $N$
- Batch Size:  $B$
- Input batch:  $X$        $X.shape = (B, N)$
- Layer input size:  $N$
- Layer output size:  $M$
- Layer weights:  $W$        $W.shape = (N, M)$
- Layer bias:  $b$        $b.shape = (M, 1)$
- Layer output:  $WX + b$        $(WX + b).shape = (B, M)$

# Applying Activation Functions

- We have output from the layer ( $WX + b$ )
- Most of activation functions apply function to the every entry in the matrix individually
- For instance, sigmoid:

$$A = \begin{pmatrix} 6 & 2 & 3 \\ 1 & 5 & 7 \end{pmatrix}$$

$$\sigma(A) = \begin{pmatrix} \sigma(6) & \sigma(2) & \sigma(3) \\ \sigma(1) & \sigma(5) & \sigma(7) \end{pmatrix}$$

- An exception – *softmax*. But is still quite straightforward

# Why Do We Care About These Matrices

- It is faster
- People wrote a lot of code for efficient matrix operations
- Graphical Processing Units can process this even more efficiently

# Matrices Recap

- We can represent our batch as a matrix  $X$  of shape  $(B, N)$
- We can represent our dense layer *weights* with a matrix  $W$  of shape  $(Inputs, Outputs)$
- We can represent our dense layer bias with a vector  $b$  of shape  $(Outputs, 1)$
- We can compute dense layer output via  $WX + b$
- And we apply non-linearity just like that  $f(WX + b)$

# Lecture Recap

- Logistic regression is a 1-layer Neural Network
- We optimize it with gradient descent
- To compute gradients for deeper models we use *backpropagation*
- We also train models in a mini-batch setting
- We can represent our neural networks via matrix multiplications

This is It For the Second Lecture